



CLOUDFLARE™

Understanding Go Memory

September 11, 2013

John Graham-Cumming

Allocation Primitives

- `new(T)`
 - Allocates memory for item with type `T`
 - Zeroes it

```
ret = runtime·mallocgc(typ->size, flag, 1, 1);
```

zeroed

- `make(T)`
 - Allocates memory for item with type `T`
 - Initializes it
 - Needed for channels, maps and slices
- Memory comes from internal heap

Two memory freeing processes

- Garbage collection
 - Determines which blocks of memory are no longer used
 - Marks areas of heap so they can be reused by your program
- Scavenging
 - Determines when parts of the heap are idle for a long time
 - Returns memory to the operation system

Garbage collection

- Controlled by the GOGC environment variable
 - And by `debug.SetGCPercentage()`

```
// Initialized from $GOGC.  GOGC=off means no gc.
//
// Next gc is after we've allocated an extra amount of
// memory proportional to the amount already in use.
// If gcpercent=100 and we're using 4M, we'll gc again
// when we get to 8M.  This keeps the gc cost in linear
// proportion to the allocation cost.  Adjusting gcpercent
// just changes the linear constant (and also the amount of
// extra memory used).
```

- Default is same as `GOGC=100`
- Can set `GOGC=off` or `debug.SetGCPercentage(-1)` (no garbage collection at all)

Scavenging

- Runs once per minute

```
// If we go two minutes without a garbage collection,  
// force one to run.  
forcegc = 2*60*1e9;
```

```
// If a span goes unused for 5 minutes after a garbage  
// collection, we hand it back to the operating system.  
limit = 5*60*1e9;
```

- Can also force return of all unused memory by calling `debug.FreeOSMemory()`


Memory Statistics

- Read with `runtime.ReadMemStats(&m)`
- The `MemStats` struct has tons of members
- Useful ones for looking at heap
 - `HeapInuse` - # bytes in the heap allocated to things
 - `HeapIdle` - # bytes in heap waiting to be used
 - `HeapSys` - # bytes obtained from OS
 - `HeapReleased` - # bytes released to OS


Test garbage making program

```
func makeBuffer() []byte {  
    return make([]byte, rand.Intn(5000000)+5000000)  
}  
  
func main() {  
    pool := make([][]byte, 20)  
  
    makes := 0  
    for {  
        b := makeBuffer()  
        makes += 1  
  
        i := rand.Intn(len(pool))  
        pool[i] = b  
  
        time.Sleep(time.Second)  
    }  
}
```

What happens



debug.FreeOSMemory()



Use a buffered channel

```
func main() {  
    pool := make([][]byte, 20)  
    idle:= make(chan []byte, 5)
```

```
  
    makes := 0  
    for {  
        var b []byte  
        select {  
        case b = <-idle:  
        default:  
            makes += 1  
            b = makeBuffer()  
        }  
    }
```

```
        i := rand.Intn(len(pool))  
        if pool[i] != nil {  
            select {  
            case idle<- pool[i]:  
                pool[i] = nil  
            default:  
            }  
        }  
  
        pool[i] = b  
  
        time.Sleep(time.Second)  
    }  
}
```

select for non-blocking receive

A buffered channel makes a simple queue

```
idle:= make(chan []byte, 5)
```

```
select {  
case b = <-idle:
```

```
default:  
    makes += 1  
    b = makeBuffer()  
}
```

Try to get from the idle queue

Idle queue empty? Make a new buffer

select for non-blocking send

A buffered channel makes a simple queue

```
idle:= make(chan []byte, 5)
```


```
select {  
case buffer <- pool[i]:  
    pool[i] = nil
```

```
default:  
}
```

Try to return buffer to the idle queue

Idle queue full?
GC will have to deal with the buffer

What happens




More realistic: 20 goroutines

```
func main() {
    pool := make([][]byte, 200)

    for i := 0; i < 10; i++ {
        go func(offset int) {
            for {
                b := makeBuffer()
                j := offset+rand.Intn(20)
                pool[j] = b

                time.Sleep(time.Millisecond * time.Duration(rand.Intn(1000)))
            }
        }(i*20)
    }
}
```

What happens




Shared across goroutines

```
func main() {
    buffer := make(chan []byte, 5)

    pool := make([][]byte, 200)
    for i := 0; i < 10; i++ {
        go func(offset int) {
            for {
                var b []byte
                select {
                    case b = <-buffer:
                    default: b = makeBuffer()
                }
                j := offset+rand.Intn(20)
                if pool[j] != nil {
                    select {
                        case buffer <- pool[j]: pool[j] = nil
                        default:
                    }
                }
                pool[j] = b
                time.Sleep(time.Millisecond * time.Duration(rand.Intn(1000)))
            }
        }(i*20)
    }
}
```



What Happens




More realistic example

- Alter code to
 - Always try to give back a random buffer from the pool
 - 50% of the time get a new one
- Should create more garbage


Idle length 5



Idle length 20



Idle length 50



Also

- This works for things other than []byte
 - Can be done with arbitrary types
 - Just need some way to reset
- There's a proposal to add something like this to the Go package library
 - sync.Cache
 - Follow TODO