



CLOUDFLARE™

Profiling Go Programs

November 6, 2013

John Graham-Cumming

The Golden Rule

- Premature optimization is the root of all evil – Hoare
- My version...
 - Don't waste programmer cycles saving the wrong CPU cycles
 - (applies to memory as well)
- Measure, measure, measure
 - Then you can optimize

Some Go Measurement Tools

- **Timing**

- **Shell `time` command**
- `time.Now()`, `time.Duration()`
- `pprof.StartCPUProfile(f)`, `pprof.StopCPUProfile()`
- `go tool pprof http://localhost:6060/debug/pprof/profile`

- **Memory**

- **Shell `ps` command**
- `runtime.ReadMemStats(&m)`
- `runtime.WriteHeapProfile(w)`
- `go tool pprof http://localhost:6060/debug/pprof/heap`

CPU PROFILING

Example: badly implemented LRU Cache

```
package lru1

import "time"

type Item struct {
    key string
    value string
    last time.Time
}

type Cache struct {
    cap int
    data map[string]*Item
}

func NewCache(cap int) (*Cache) {
    return &Cache{cap, make(map[string]*Item)}
}
```

Example: badly implemented LRU Cache

```
func (c *Cache) makeSpace() {
    old := &Item{last: time.Now()}
    var key string
    for k, v := range c.data {
        if v.last.Before(old.last) {
            old = v
            key = k
        }
    }

    delete(c.data, key)
}

func (c *Cache) Put(key, value string) {
    if len(c.data) == c.cap {
        c.makeSpace()
    }

    c.data[key] = &Item{value, time.Now()}
}
```

Example: badly implemented LRU Cache

```
func (c *Cache) Get(key string) (*Item) {  
    if c.data[key] != nil {  
        c.data[key].last = time.Now()  
    }  
  
    return c.data[key]  
}
```

Bad LRU Cache

- Use it to cache relationship between email addresses and their domain's MX record
- Feed in 10M email addresses in arrival order; cache 10,000 MX records

```
% time ./lrutest1 < top10M
9929964 total 2565368 misses
82.39s user 25.22s system 99% cpu 1:47.61 total
```

- So 1m48s

Let's profile it!

- Use the pprof profiler

```
func main() {  
    f, _ := os.Create("lrutest1.cpubprofile")  
    pprof.StartCPUProfile(f)  
    defer pprof.StopCPUProfile()  
  
    ...  
}
```

- **Creates** lrutest1.cpubprofile
- `go tool pprof lrutest1 lrutest1.cpubprofile`

pprof command-line

```
% go tool pprof lrutest1 lrutest1.cpuprofile
```

```
Welcome to pprof!  For help, type 'help'.
```

```
(pprof) top
```

```
Total: 10440 samples
```

3826	36.6%	36.6%	3826	36.6%	hash_next
2252	21.6%	58.2%	9004	86.2%	lru1.(*Cache).makeSpace
2130	20.4%	78.6%	2920	28.0%	runtime.mapiter2
623	6.0%	84.6%	623	6.0%	runtime.memcopy64
248	2.4%	87.0%	3916	37.5%	runtime.mapiternext
242	2.3%	89.3%	777	7.4%	strings.genSplit
168	1.6%	90.9%	168	1.6%	runtime.strcopy
151	1.4%	92.3%	151	1.4%	strings.Count
64	0.6%	93.0%	195	1.9%	scanblock
62	0.6%	93.5%	390	3.7%	runtime.mallocgc

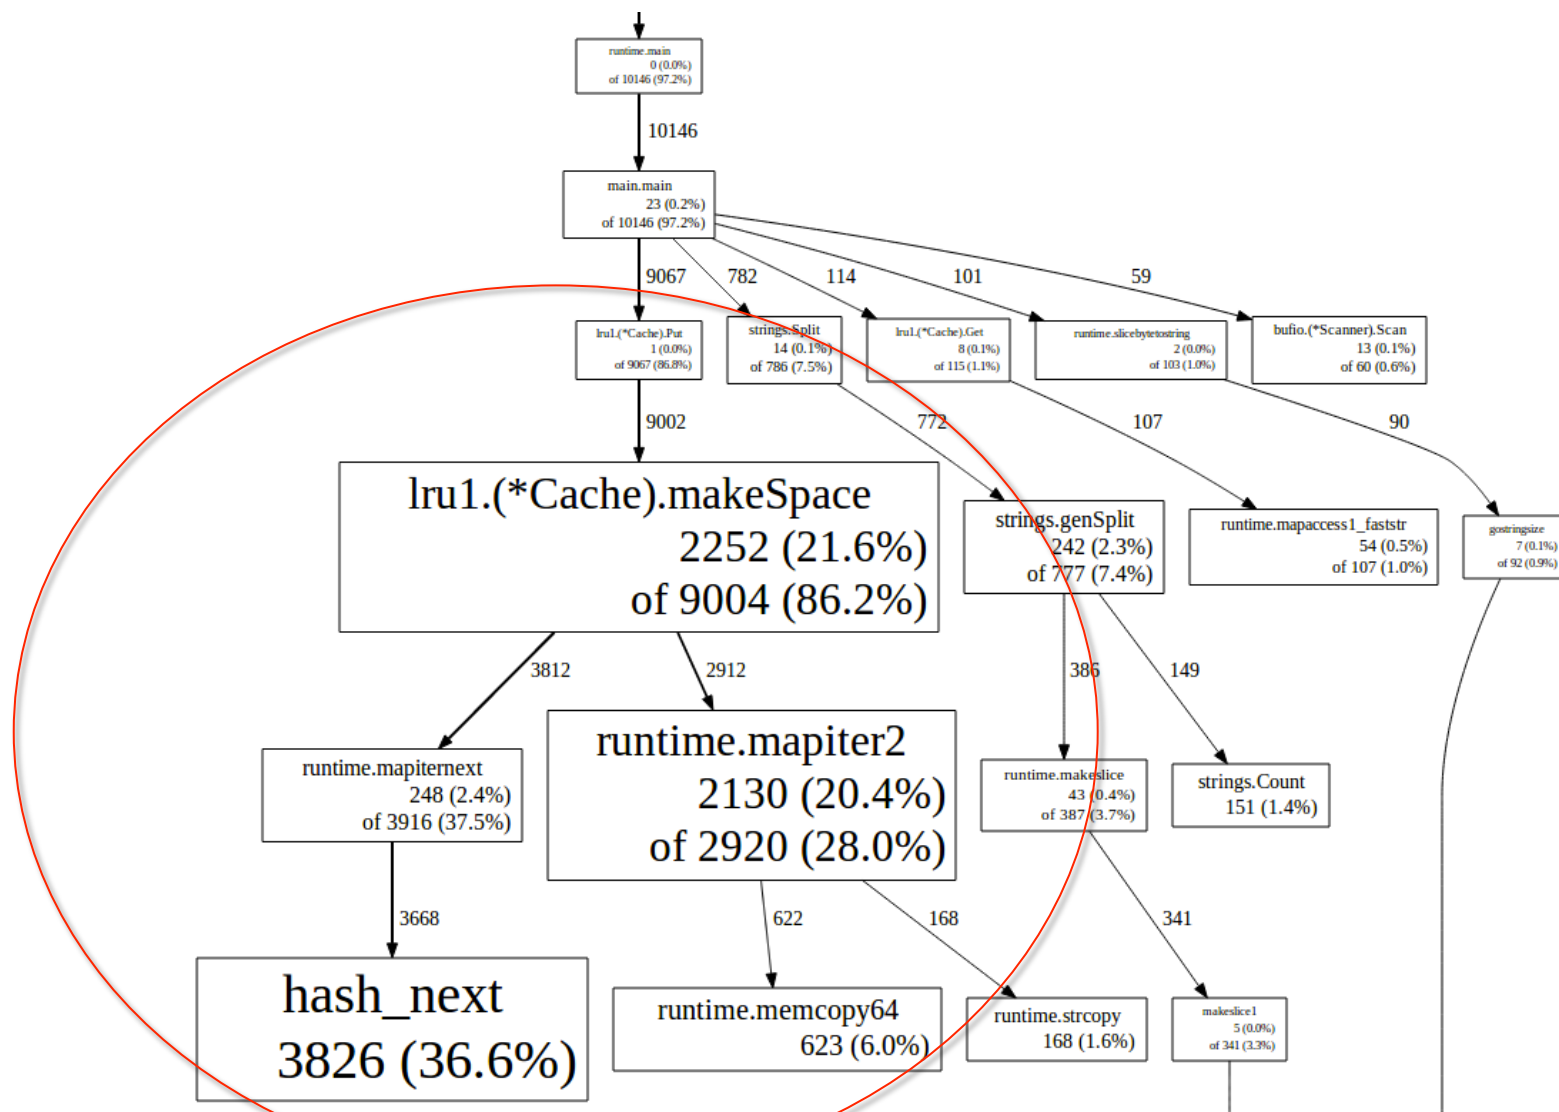
pprof command-line

```
(pprof) top -cum
```

```
Total: 10440 samples
```

0	0.0%	0.0%	10146	97.2%	gosched0
23	0.2%	0.2%	10146	97.2%	main.main
0	0.0%	0.2%	10146	97.2%	runtime.main
1	0.0%	0.2%	9067	86.8%	lru1.(*Cache).Put
2252	21.6%	21.8%	9004	86.2%	lru1.(*Cache).makeSpace
248	2.4%	24.2%	3916	37.5%	runtime.mapiternext
3826	36.6%	60.8%	3826	36.6%	hash_next
2130	20.4%	81.2%	2920	28.0%	runtime.mapiter2
14	0.1%	81.4%	786	7.5%	strings.Split
242	2.3%	83.7%	777	7.4%	strings.genSplit

pprof web view



Live profiling

- `net/http/pprof`

```
import _ "net/http/pprof"
```

```
go func() {  
    log.Println(http.ListenAndServe("127.0.0.1:6161", nil))  
}()
```

- `/pprof/debug/heap`
- `/pprof/debug/profile`

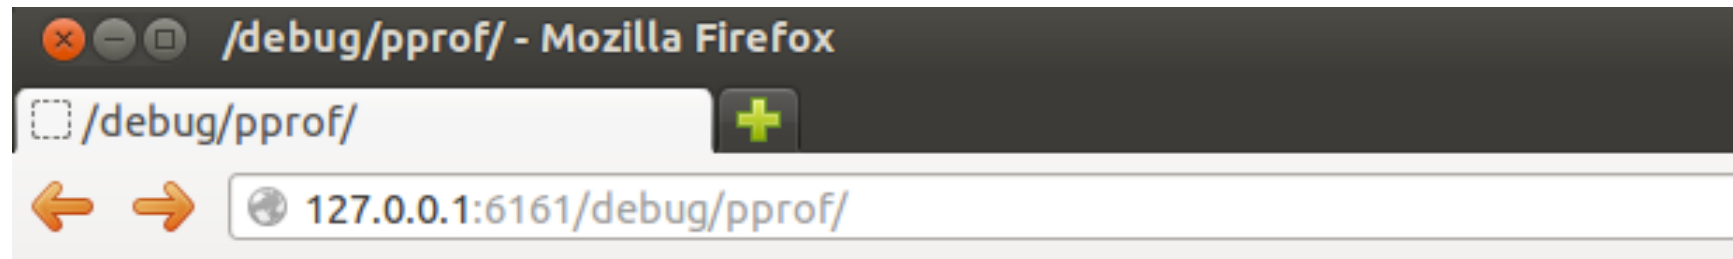
Live command-line profiling

```
% go tool pprof http://127.0.0.1:6161/debug/pprof/profile
Read http://127.0.0.1:6161/debug/pprof/symbol
Be patient...
Wrote profile to [...]
Welcome to pprof! For help, type 'help'.
(pprof) top
Total: 2948 samples
```

1146	38.9%	38.9%	1146	38.9%	hash_next
642	21.8%	60.7%	2618	88.8%	lru1.(*Cache).makeSpace
582	19.7%	80.4%	806	27.3%	runtime.mapiter2
176	6.0%	86.4%	176	6.0%	runtime.memcopy64
86	2.9%	89.3%	1194	40.5%	runtime.mapiternext
51	1.7%	91.0%	176	6.0%	strings.genSplit
48	1.6%	92.6%	48	1.6%	runtime.strcopy
43	1.5%	94.1%	43	1.5%	runtime.futex
43	1.5%	95.6%	43	1.5%	strings.Count
14	0.5%	96.0%	85	2.9%	runtime.makeslice

```
(pprof)
```

Live browser profiling



/debug/pprof/

profiles:

0 [block](#)

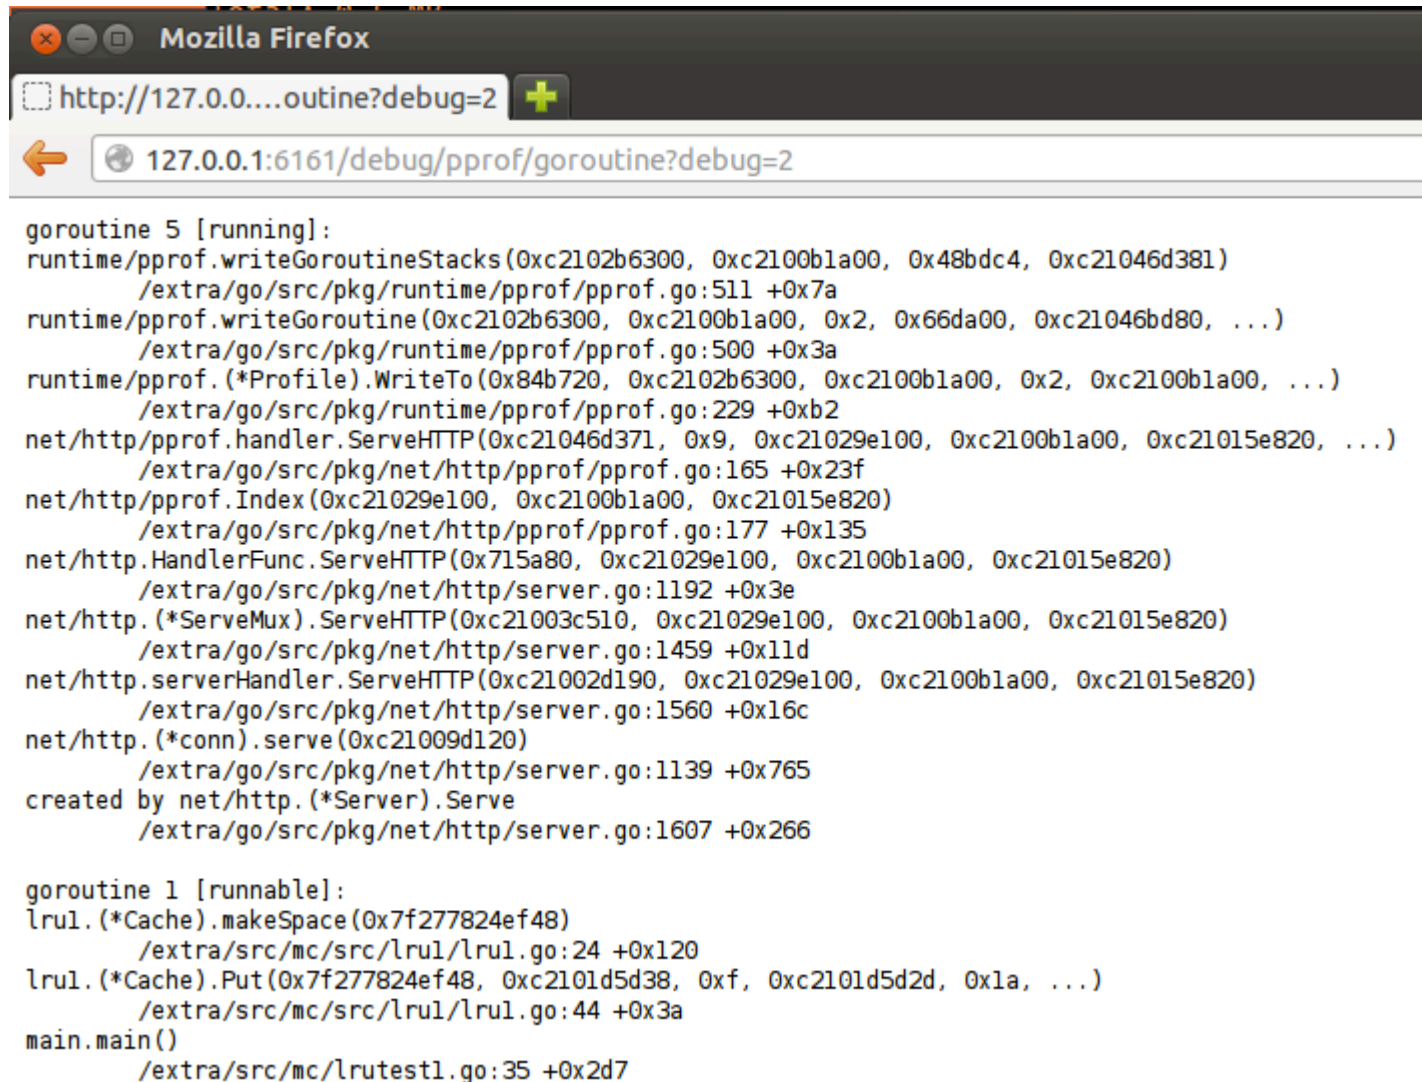
5 [goroutine](#)

4 [heap](#)

13 [threadcreate](#)

[full goroutine stack dump](#)

Full goroutine stack dump



```
goroutine 5 [running]:
runtime/pprof.writeGoroutineStacks(0xc2102b6300, 0xc2100b1a00, 0x48bdc4, 0xc21046d381)
    /extra/go/src/pkg/runtime/pprof/pprof.go:511 +0x7a
runtime/pprof.writeGoroutine(0xc2102b6300, 0xc2100b1a00, 0x2, 0x66da00, 0xc21046bd80, ...)
    /extra/go/src/pkg/runtime/pprof/pprof.go:500 +0x3a
runtime/pprof.(*Profile).WriteTo(0x84b720, 0xc2102b6300, 0xc2100b1a00, 0x2, 0xc2100b1a00, ...)
    /extra/go/src/pkg/runtime/pprof/pprof.go:229 +0xb2
net/http/pprof.handler.ServeHTTP(0xc21046d371, 0x9, 0xc21029e100, 0xc2100b1a00, 0xc21015e820, ...)
    /extra/go/src/pkg/net/http/pprof/pprof.go:165 +0x23f
net/http/pprof.Index(0xc21029e100, 0xc2100b1a00, 0xc21015e820)
    /extra/go/src/pkg/net/http/pprof/pprof.go:177 +0x135
net/http.HandlerFunc.ServeHTTP(0x715a80, 0xc21029e100, 0xc2100b1a00, 0xc21015e820)
    /extra/go/src/pkg/net/http/server.go:1192 +0x3e
net/http.(*ServeMux).ServeHTTP(0xc21003c510, 0xc21029e100, 0xc2100b1a00, 0xc21015e820)
    /extra/go/src/pkg/net/http/server.go:1459 +0x11d
net/http.serverHandler.ServeHTTP(0xc21002d190, 0xc21029e100, 0xc2100b1a00, 0xc21015e820)
    /extra/go/src/pkg/net/http/server.go:1560 +0x16c
net/http.(*conn).serve(0xc21009d120)
    /extra/go/src/pkg/net/http/server.go:1139 +0x765
created by net/http.(*Server).Serve
    /extra/go/src/pkg/net/http/server.go:1607 +0x266

goroutine 1 [runnable]:
lrul.(*Cache).makeSpace(0x7f277824ef48)
    /extra/src/mc/src/lrul/lrul.go:24 +0x120
lrul.(*Cache).Put(0x7f277824ef48, 0xc2101d5d38, 0xf, 0xc2101d5d2d, 0x1a, ...)
    /extra/src/mc/src/lrul/lrul.go:44 +0x3a
main.main()
    /extra/src/mc/lrutest1.go:35 +0x2d7
```


Running goroutines

```
Mozilla Firefox
http://127.0.0.1:6161/debug/pprof/goroutine?debug=1

goroutine profile: total 4
1 @ 0x5004ae 0x5002b7 0x4fd2e2 0x44d12f 0x44d275 0x43873e 0x439a4d 0x43a21c 0x4383e5 0x417900
# 0x5004ae runtime/pprof.writeRuntimeProfile+0x9e /extra/go/src/pkg/runtime/pprof/pprof.go:540
# 0x5002b7 runtime/pprof.writeGoroutine+0x87 /extra/go/src/pkg/runtime/pprof/pprof.go:502
# 0x4fd2e2 runtime/pprof.(*Profile).WriteTo+0xb2 /extra/go/src/pkg/runtime/pprof/pprof.go:229
# 0x44d12f net/http/pprof.handler.ServeHTTP+0x23f /extra/go/src/pkg/net/http/pprof/pprof.go:165
# 0x44d275 net/http/pprof.Index+0x135 /extra/go/src/pkg/net/http/pprof/pprof.go:177
# 0x43873e net/http.HandlerFunc.ServeHTTP+0x3e /extra/go/src/pkg/net/http/server.go:1192
# 0x439a4d net/http.(*ServeMux).ServeHTTP+0x11d /extra/go/src/pkg/net/http/server.go:1459
# 0x43a21c net/http.serverHandler.ServeHTTP+0x16c /extra/go/src/pkg/net/http/server.go:1560
# 0x4383e5 net/http.(*conn).serve+0x765 /extra/go/src/pkg/net/http/server.go:1139

1 @ 0x417855 0x409df6 0x45f3f0 0x45f58a 0x400ed7 0x415912 0x417900
# 0x45f3f0 lru1.(*Cache).makeSpace+0x120 /extra/src/mc/src/lru1/lru1.go:24
# 0x45f58a lru1.(*Cache).Put+0x3a /extra/src/mc/src/lru1/lru1.go:44
# 0x400ed7 main.main+0x2d7 /extra/src/mc/lrutest1.go:35
# 0x415912 runtime.main+0x92 /extra/go/src/pkg/runtime/proc.c:181

1 @ 0x417e6e 0x41258b 0x417900
# 0x417e6e runtime.entersyscallblock+0x16e /extra/go/src/pkg/runtime/proc.c:1342
# 0x41258b runtime.MHeap_Scavenger+0xeb /extra/go/src/pkg/runtime/mheap.c:471

1 @ 0x417744 0x4219af 0x421472 0x4c7371 0x4cald1 0x4db5a5 0x4db675 0x43a395 0x43a2ee 0x43a5e5 0x40100e 0x417900
# 0x4219af netpollblock+0x9f /extra/go/src/pkg/runtime/znetpoll_linux_amd64.c:255
# 0x421472 net.runtime_pollWait+0x82 /extra/go/src/pkg/runtime/znetpoll_linux_amd64.c:118
# 0x4c7371 net.(*pollDesc).WaitRead+0x31 /extra/go/src/pkg/net/fd_poll_runtime.go:70
# 0x4cald1 net.(*netFD).accept+0x2c1 /extra/go/src/pkg/net/fd_unix.go:390
# 0x4db5a5 net.(*TCPLListener).AcceptTCP+0x45 /extra/go/src/pkg/net/tcpsock_posix.go:229
# 0x4db675 net.(*TCPLListener).Accept+0x25 /extra/go/src/pkg/net/tcpsock_posix.go:239
# 0x43a395 net/http.(*Server).Serve+0x85 /extra/go/src/pkg/net/http/server.go:1585
# 0x43a2ee net/http.(*Server).ListenAndServe+0x9e /extra/go/src/pkg/net/http/server.go:1575
# 0x43a5e5 net/http.ListenAndServe+0x65 /extra/go/src/pkg/net/http/server.go:1640
# 0x40100e main.func·001+0x3e /extra/src/mc/lrutest1.go:18
```

Better LRU implementation

```
package lru2

import "container/list"

type Item struct {
    key string
    value string
}

type Cache struct {
    cap int
    data map[string]*list.Element
    l *list.List
}

func NewCache(cap int) (*Cache) {
    return &Cache{cap, make(map[string]*list.Element),
        list.New()}
}
```

Better LRU implementation

```
func (c *Cache) Get(key string) (*Item) {
    if c.data[key] != nil {
        c.l.MoveToFront(c.data[key])
        return c.data[key].Value.(*Item)
    }

    return nil
}

func (c *Cache) Put(key, value string) {
    if len(c.data) == c.cap {
        delete(c.data, c.l.Back().Value.(*Item).key)
        c.l.Remove(c.l.Back())
    }

    c.data[key] = c.l.PushFront(&Item{key, value})
}
```

Better LRU Cache

- Use it to cache relationship between email addresses and their domain's MX record
- Feed in 10M email addresses in arrival order; cache 10,000 MX records

```
% time ./lrutest2 < top10M
9929964 total 2565368 misses
12.19s user 1.14s system 105% cpu 12.652 total
```

- So 12.3s

pprof command-line

```
% go tool pprof lrutest2 lrutest2.cpuprofile
```

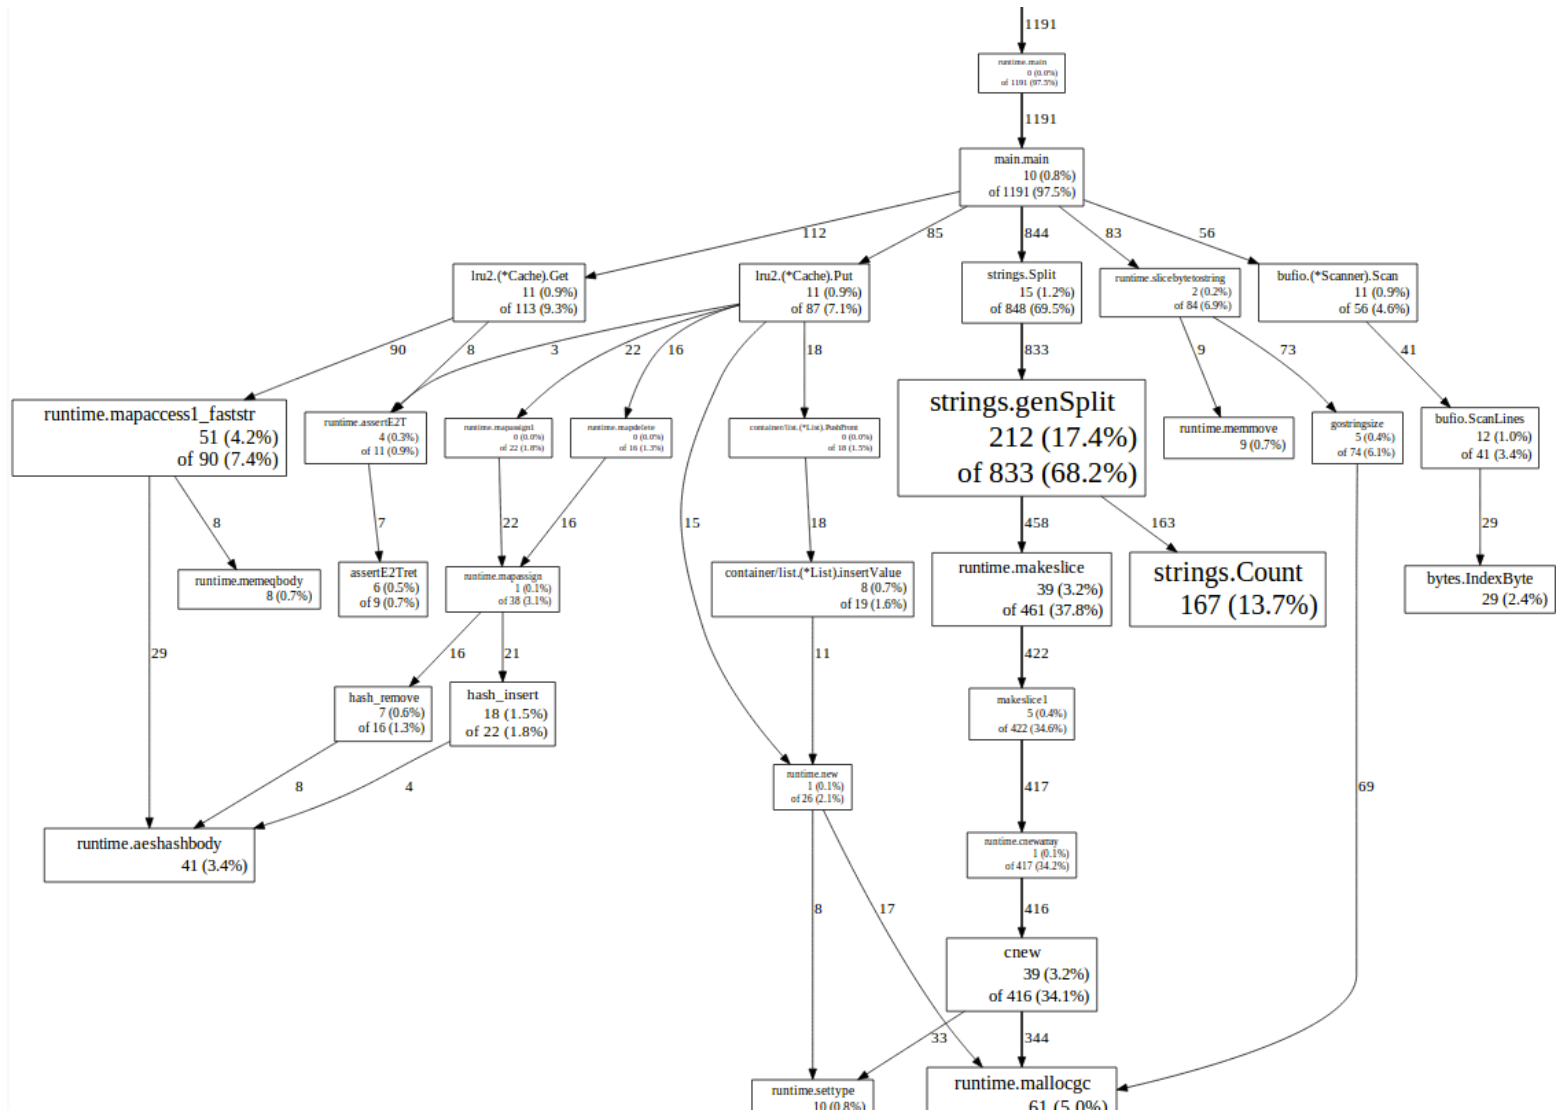
```
Welcome to pprof!  For help, type 'help'.
```

```
(pprof) top
```

```
Total: 1221 samples
```

212	17.4%	17.4%	833	68.2%	strings.genSplit
167	13.7%	31.0%	167	13.7%	strings.Count
74	6.1%	37.1%	77	6.3%	sweepspan
71	5.8%	42.9%	206	16.9%	scanblock
61	5.0%	47.9%	61	5.0%	markonly
61	5.0%	52.9%	435	35.6%	runtime.mallocgc
51	4.2%	57.1%	51	4.2%	flushptrbuf
51	4.2%	61.3%	90	7.4%	runtime.mapaccess1_faststr
49	4.0%	65.3%	49	4.0%	runtime.futex
41	3.4%	68.6%	41	3.4%	runtime.aeshashbody

pprof web view



Let's try random eviction

```
package lru3

import "math/rand"
import "fmt"

type Item struct {
    key string
    value string
}

type Cache struct {
    cap int
    data map[string]*Item
    keys []string
}

func NewCache(cap int) (*Cache) {
    return &Cache{cap, make(map[string]*Item),
        make([]string, 0, cap)}
}
```

Let's try random eviction

```
func (c *Cache) Get(key string) (*Item) {  
    return c.data[key]  
}  
  
func (c *Cache) Put(key, value string) {  
    if len(c.keys) == c.cap {  
        evict := rand.Intn(c.cap)  
        delete(c.data, c.keys[evict])  
        c.keys = append(c.keys[:evict],  
            c.keys[evict+1:]...)  
    }  
  
    c.keys = append(c.keys, key)  
  
    c.data[key] = &Item{key, value}  
}
```


Random Eviction

- Same speed for a different algorithm

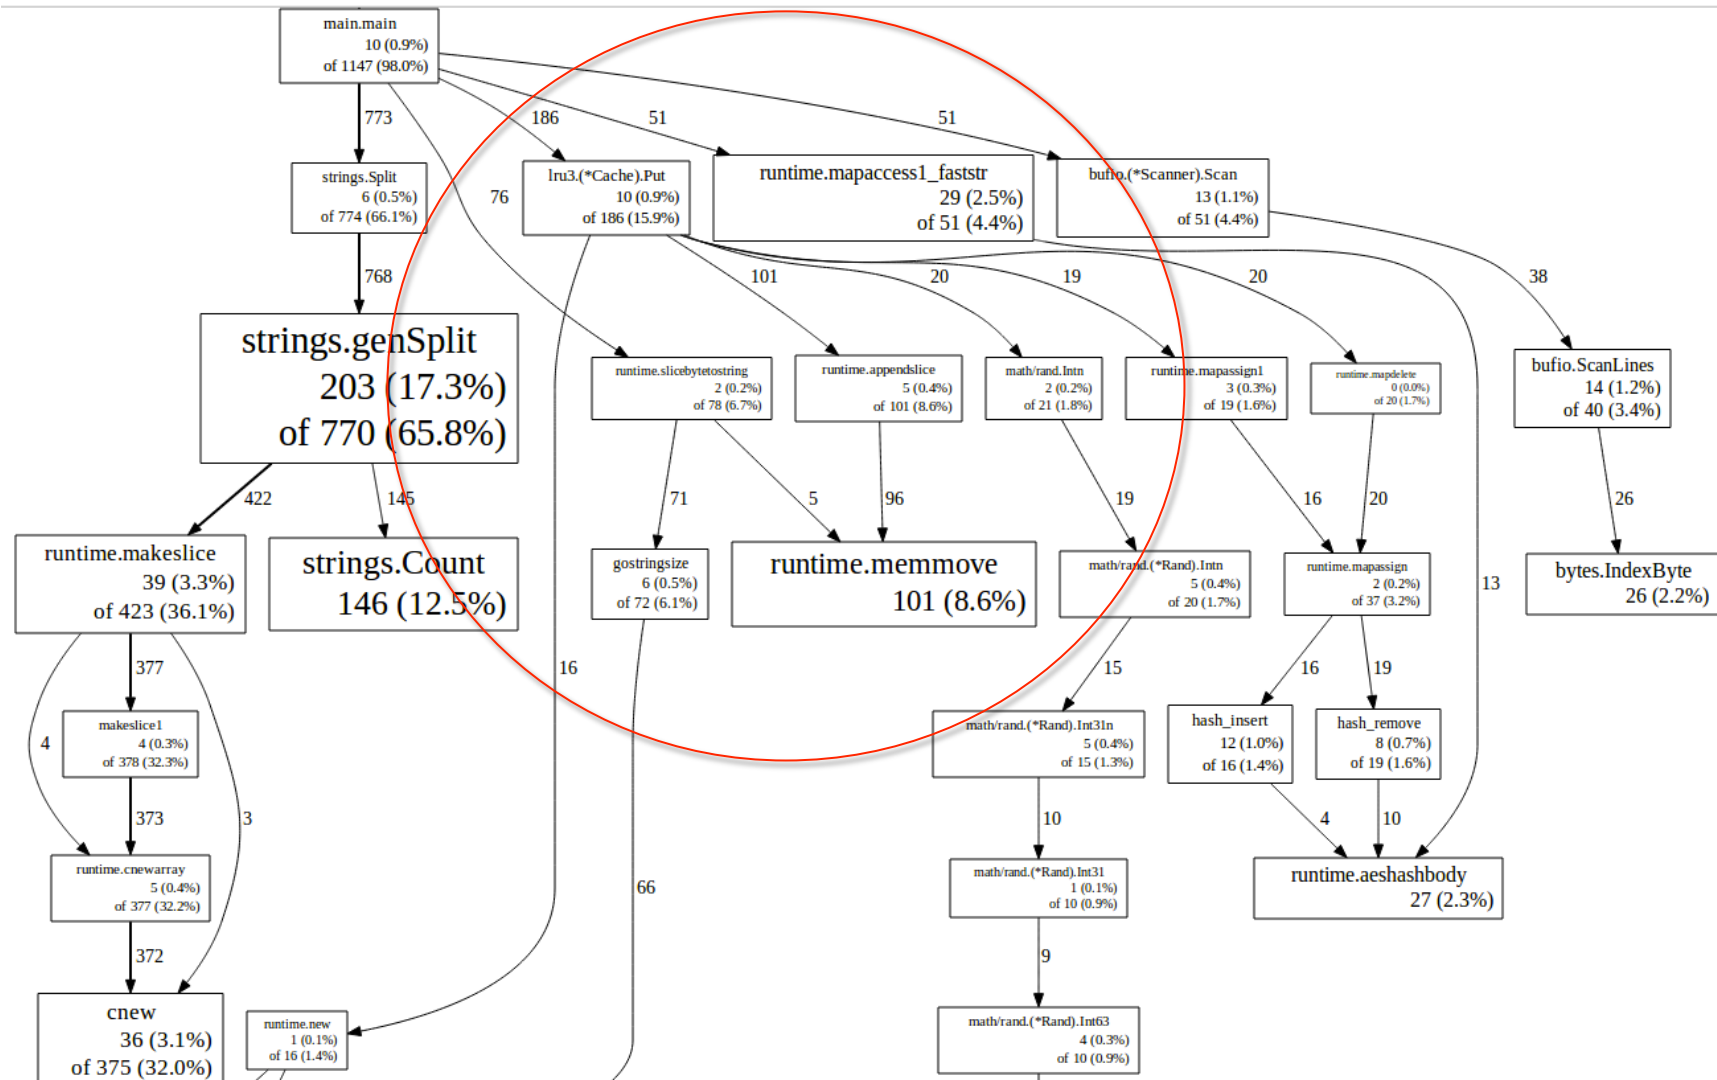
```
% time ./lrutest3 < top10M
9929964 total 2820060 misses
11.76s user 1.02s system 105% cpu 12.130 total
```

- So 12.1s

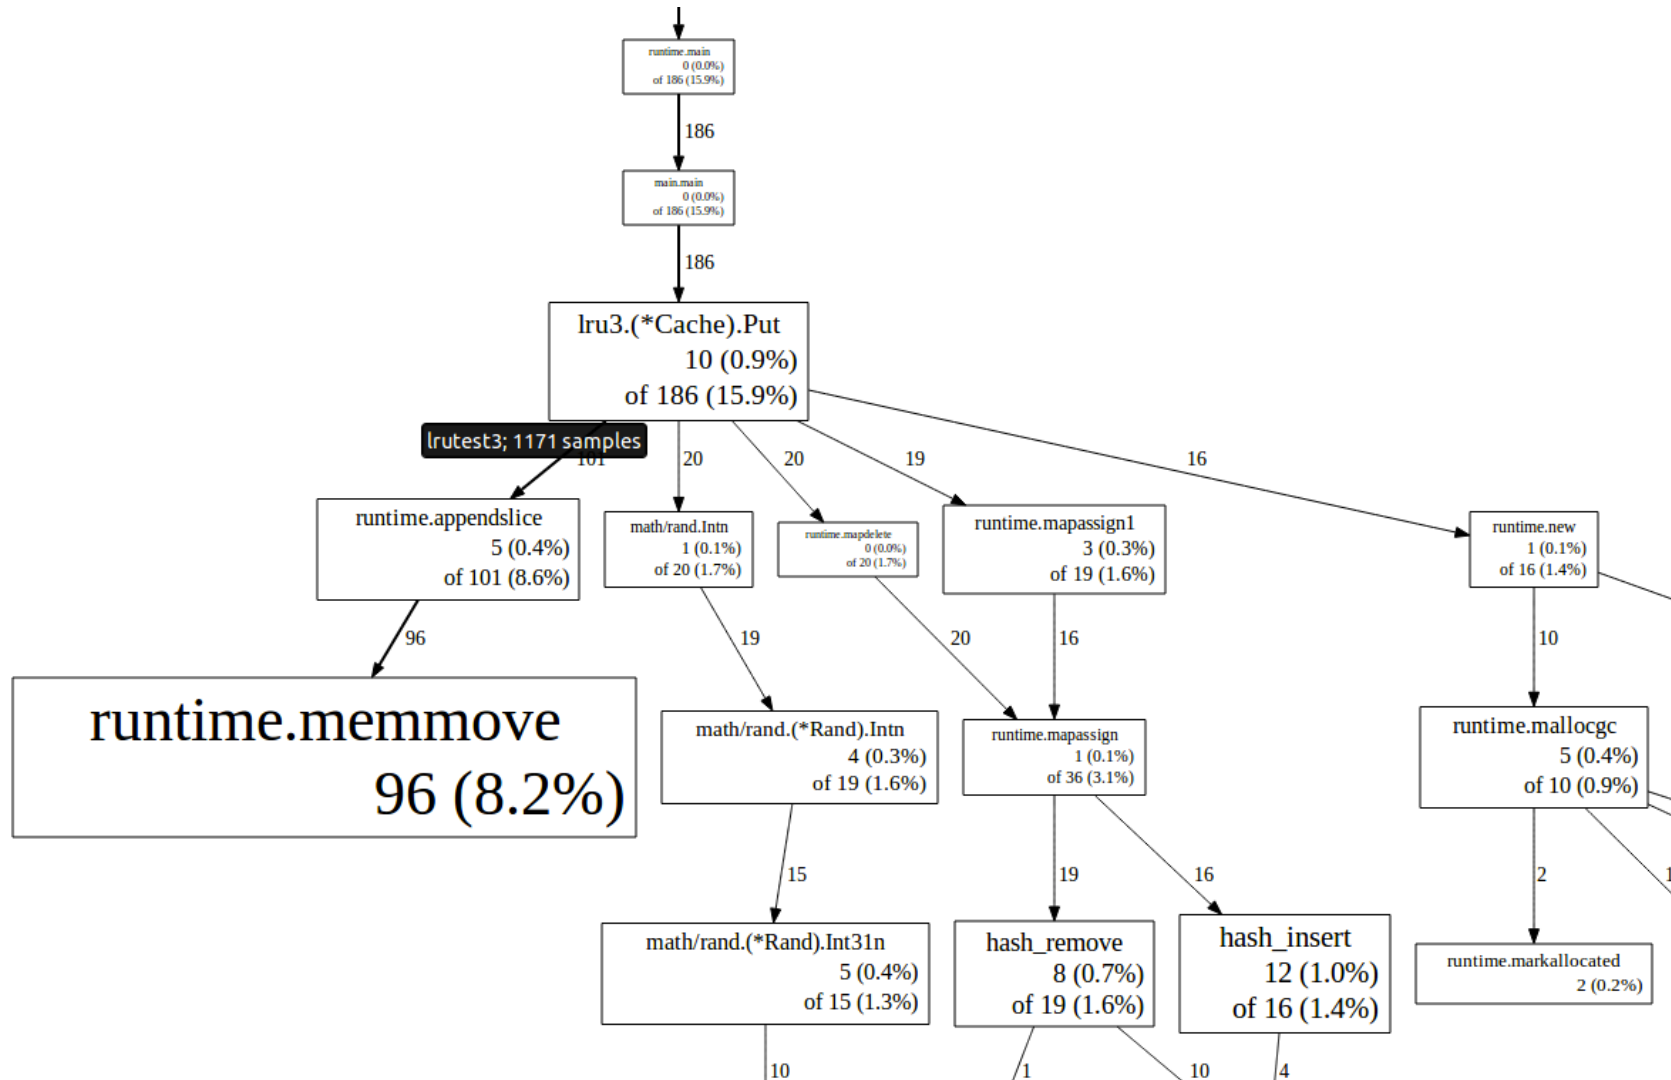
```
% go tool pprof lrutest3 lrutest3.cpubprofile
(pprof) top
Total: 1171 samples
```

203	17.3%	17.3%	770	65.8%	strings.genSplit
146	12.5%	29.8%	146	12.5%	strings.Count
101	8.6%	38.4%	101	8.6%	runtime.memmove
70	6.0%	44.4%	77	6.6%	sweepspan
61	5.2%	49.6%	380	32.5%	runtime.mallocgc
60	5.1%	54.7%	146	12.5%	scanblock
54	4.6%	59.4%	54	4.6%	markonly

pprof web view



web Put



Eliminate slice operation

```
package lru4

import "math/rand"

type Item struct {
    key string
    value string
}

type Cache struct {
    cap int
    data map[string]*Item
    keys []string
}

func NewCache(cap int) (*Cache) {
    return &Cache{cap, make(map[string]*Item),
        make([]string, cap)}
}
```

Eliminate slice operation

```
func (c *Cache) Get(key string) (*Item) {  
    return c.data[key]  
}  
  
func (c *Cache) Put(key, value string) {  
    slot := len(c.data)  
    if len(c.data) == c.cap {  
        slot = rand.Intn(c.cap)  
        delete(c.data, c.keys[slot])  
    }  
  
    c.keys[slot] = key  
    c.data[key] = &Item{key, value}  
}
```

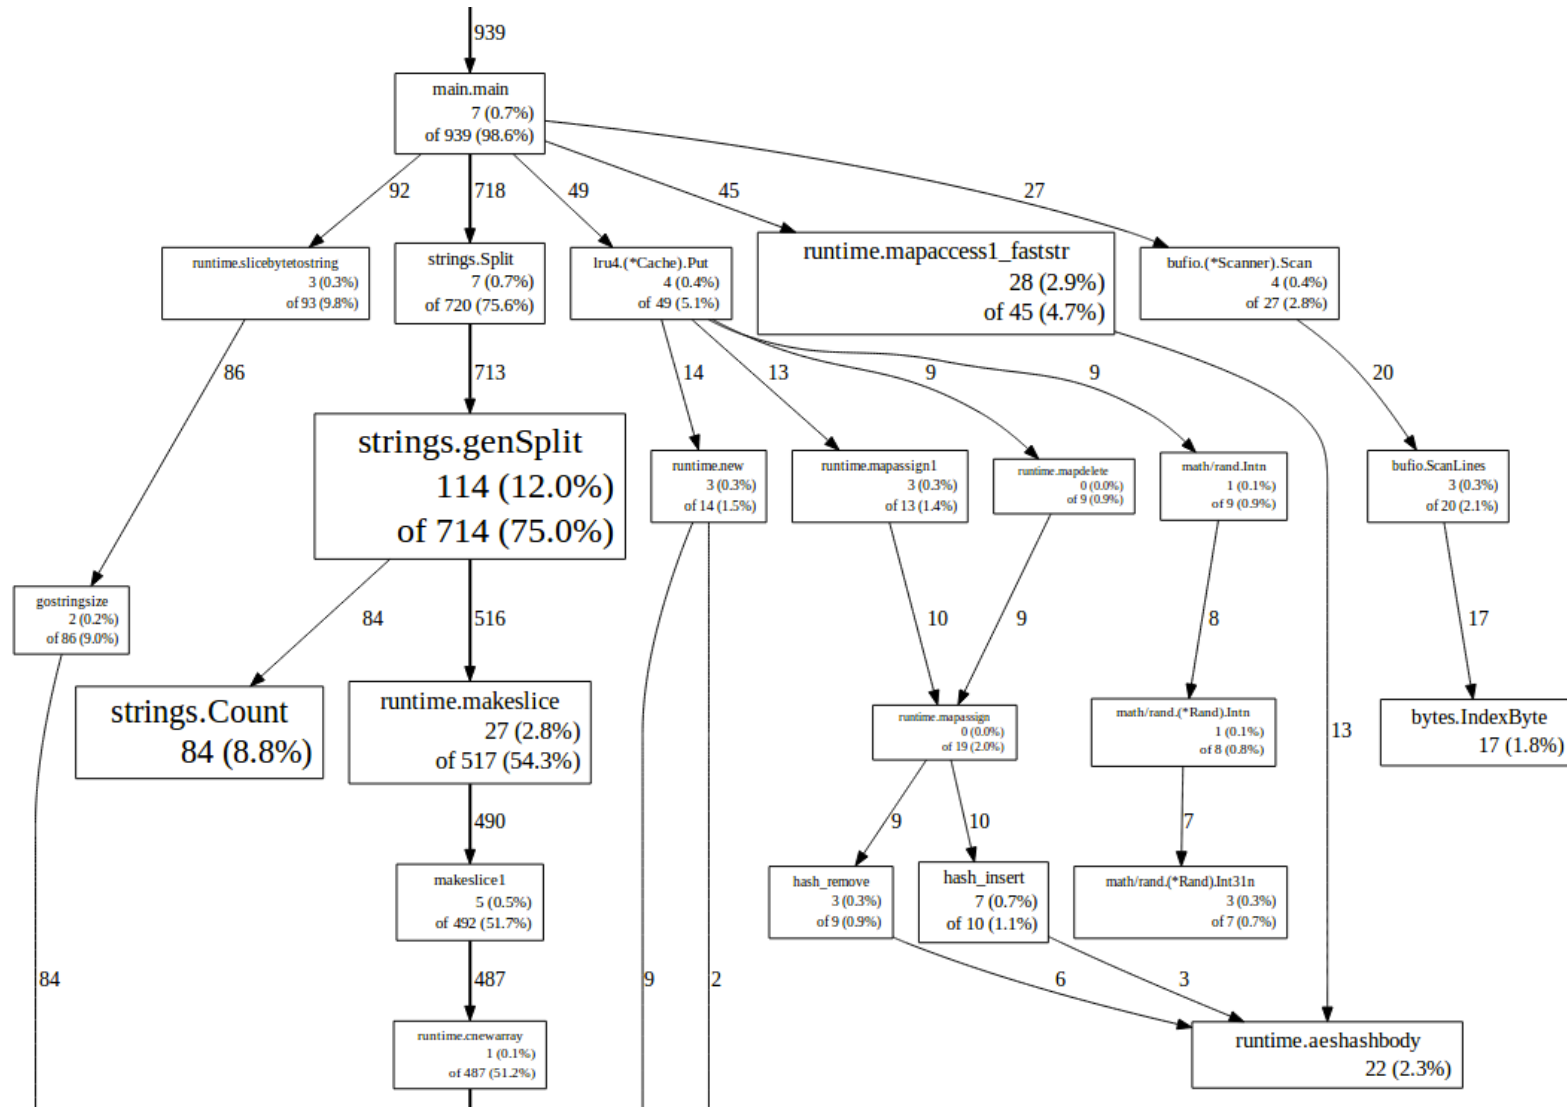
Eliminate slice operation

- Slightly faster

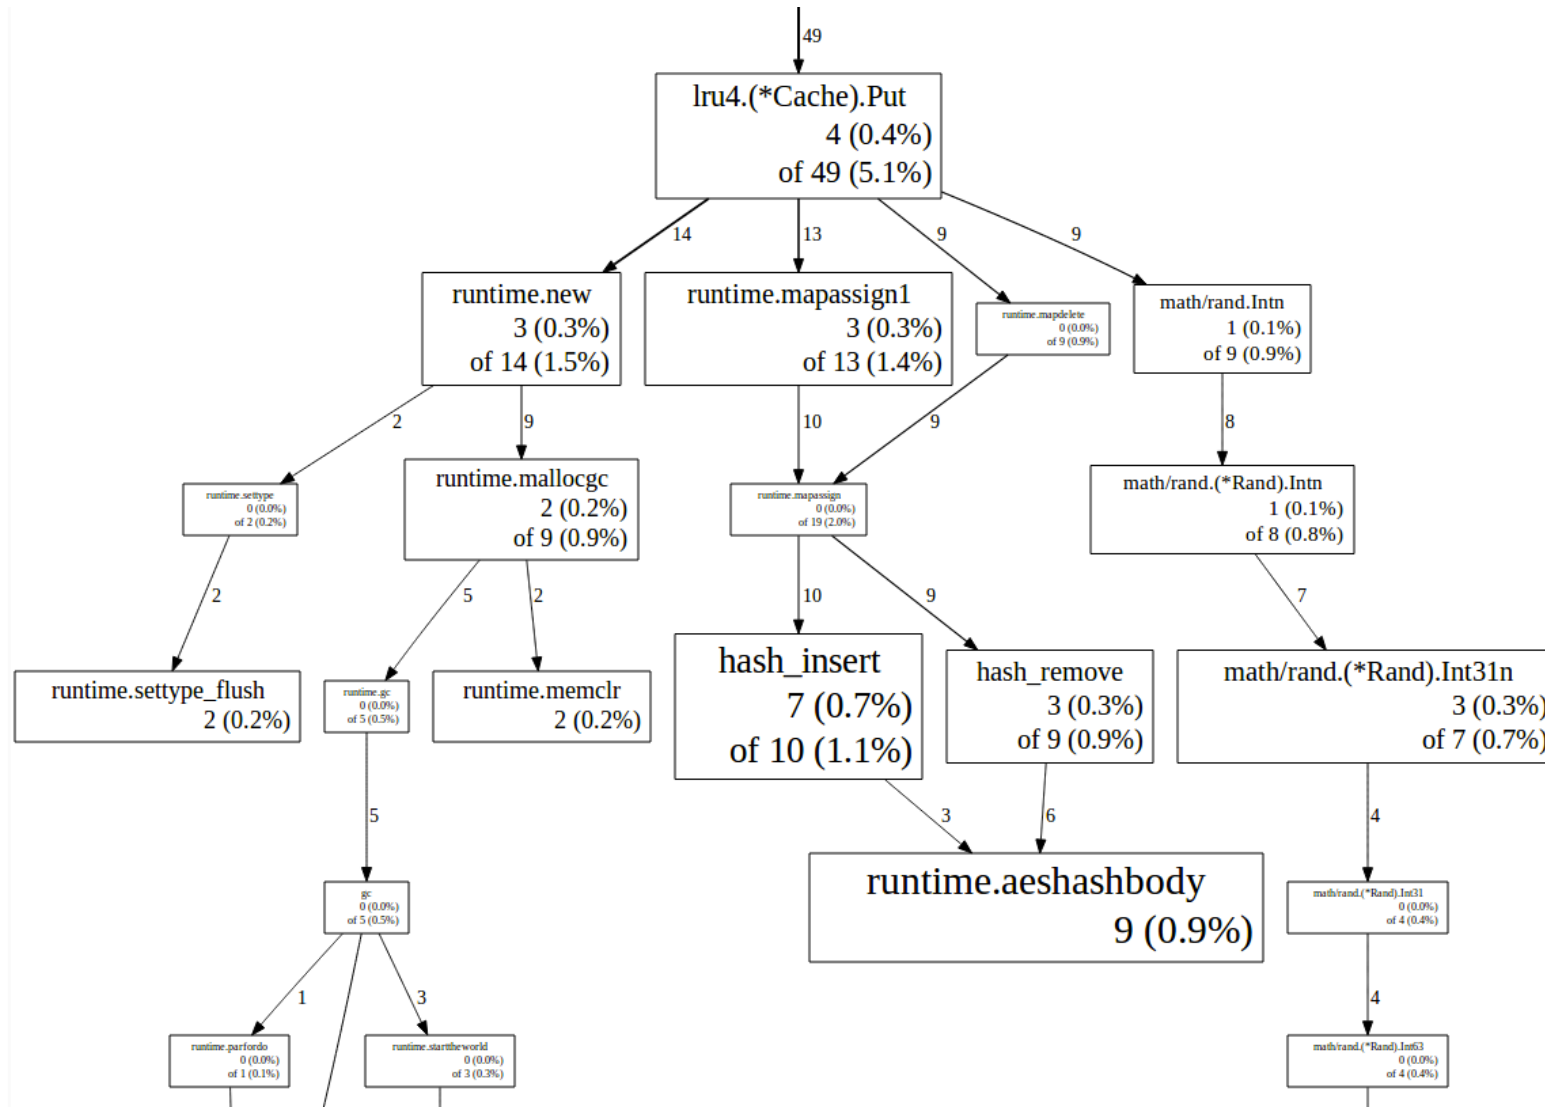
```
% time ./lrutest4 < top10M
9929964 total 2819425 misses
6.51s user 4.83s system 105% cpu 10.787 total
```

- So 10.8s

Slice operations now gone



web Put



Also...

- What's blocking on synchronization primitives?
- What's causing the creation of OS threads?

MEMORY RECYCLING

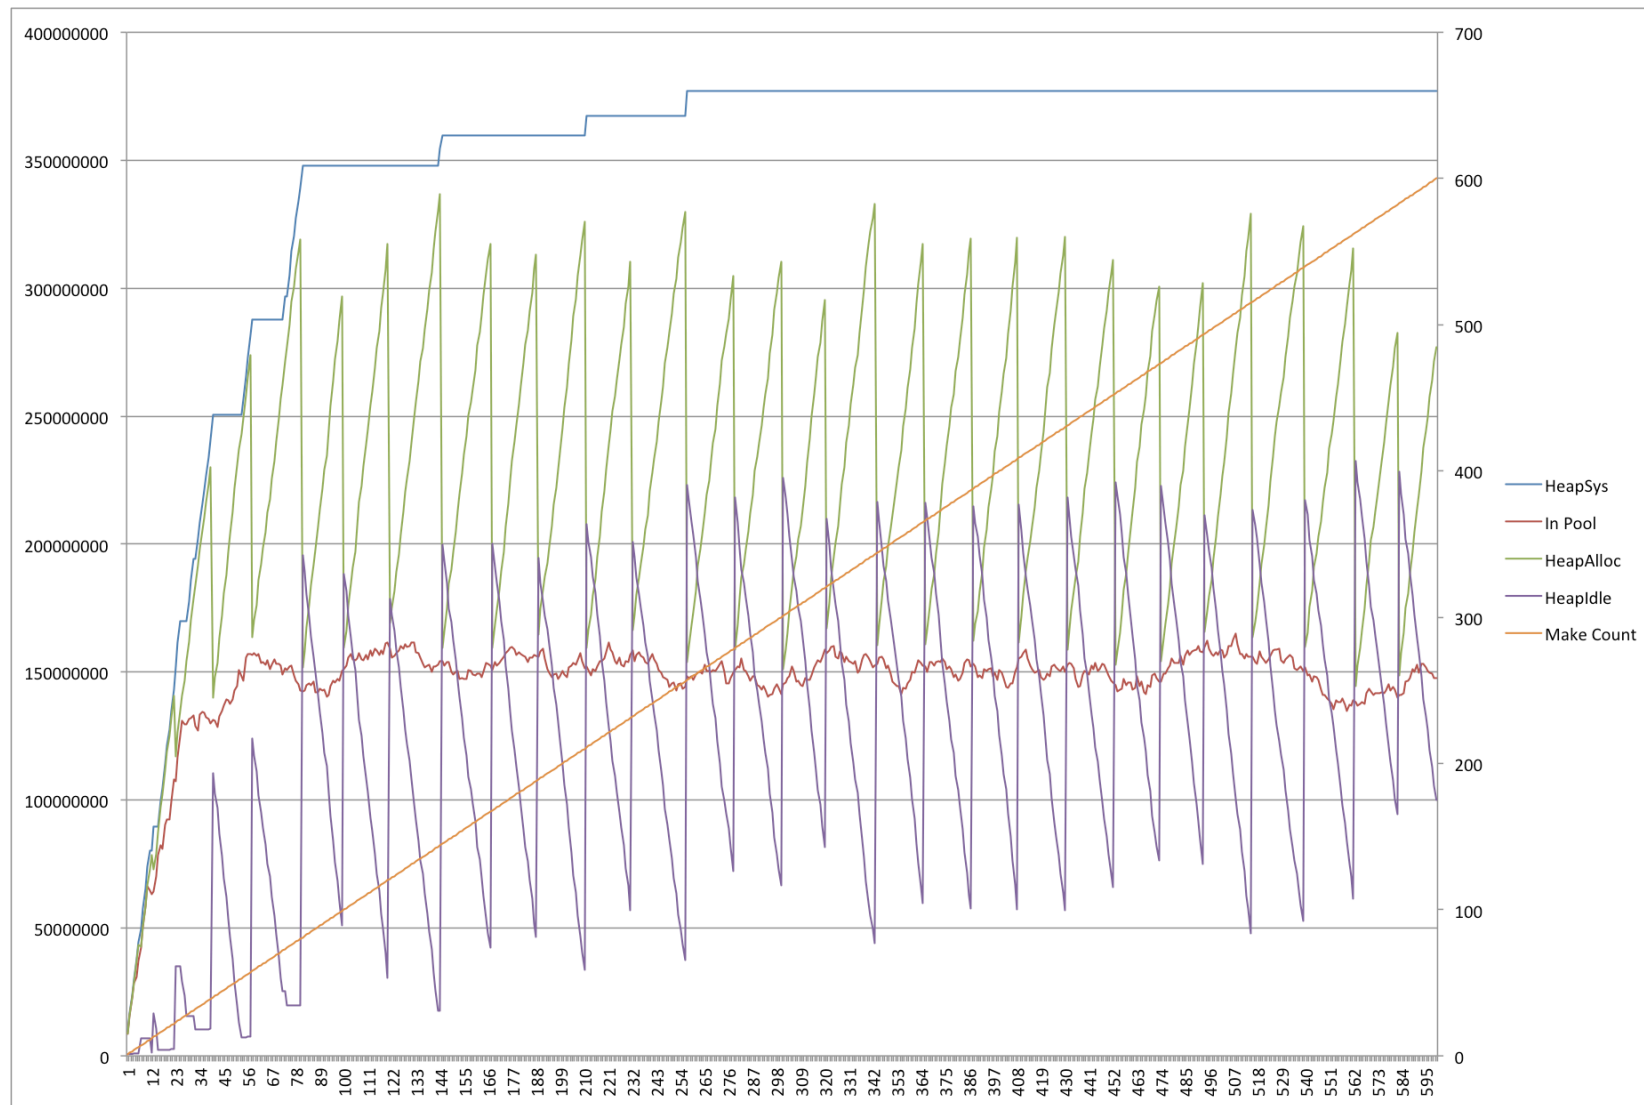
Memory Statistics

- Read with `runtime.ReadMemStats(&m)`
- The `MemStats` struct has tons of members
- Useful ones for looking at heap
 - `HeapInuse` - # bytes in the heap allocated to things
 - `HeapIdle` - # bytes in heap waiting to be used
 - `HeapSys` - # bytes obtained from OS
 - `HeapReleased` - # bytes released to OS


Test garbage making program

```
func makeBuffer() []byte {  
    return make([]byte, rand.Intn(5000000)+5000000)  
}  
  
func main() {  
    pool := make([][]byte, 20)  
  
    makes := 0  
    for {  
        b := makeBuffer()  
        makes += 1  
  
        i := rand.Intn(len(pool))  
        pool[i] = b  
  
        time.Sleep(time.Second)  
    }  
}
```

What happens



debug.FreeOSMemory()



Use a buffered channel

```
func main() {  
    pool := make([][]byte, 20)  
    idle:= make(chan []byte, 5)
```

```
  
    makes := 0  
    for {  
        var b []byte  
        select {  
        case b = <-idle:  
        default:  
            makes += 1  
            b = makeBuffer()  
        }  
    }
```

```
        i := rand.Intn(len(pool))  
        if pool[i] != nil {  
            select {  
            case idle<- pool[i]:  
                pool[i] = nil  
            default:  
            }  
        }  
  
        pool[i] = b  
  
        time.Sleep(time.Second)  
    }  
}
```

select for non-blocking receive

A buffered channel makes a simple queue

```
idle:= make(chan []byte, 5)
```

```
select {  
case b = <-idle:
```

```
default:  
    makes += 1  
    b = makeBuffer()  
}
```

Try to get from the idle queue

Idle queue empty? Make a new buffer

select for non-blocking send

A buffered channel makes a simple queue

```
idle:= make(chan []byte, 5)
```


```
select {  
case buffer <- pool[i]:  
    pool[i] = nil
```

```
default:  
}
```

Try to return buffer to the idle queue

Idle queue full?
GC will have to deal with the buffer

What happens




More realistic: 20 goroutines

```
func main() {
    pool := make([][]byte, 200)

    for i := 0; i < 10; i++ {
        go func(offset int) {
            for {
                b := makeBuffer()
                j := offset+rand.Intn(20)
                pool[j] = b

                time.Sleep(time.Millisecond * time.Duration(rand.Intn(1000)))
            }
        }(i*20)
    }
}
```

What happens




Shared across goroutines

```
func main() {
    buffer := make(chan []byte, 5)

    pool := make([][]byte, 200)
    for i := 0; i < 10; i++ {
        go func(offset int) {
            for {
                var b []byte
                select {
                    case b = <-buffer:
                    default: b = makeBuffer()
                }
                j := offset+rand.Intn(20)
                if pool[j] != nil {
                    select {
                        case buffer <- pool[j]: pool[j] = nil
                        default:
                    }
                }
                pool[j] = b
                time.Sleep(time.Millisecond * time.Duration(rand.Intn(1000)))
            }
        }(i*20)
    }
}
```


What Happens




More realistic example

- Alter code to
 - Always try to give back a random buffer from the pool
 - 50% of the time get a new one
- Should create more garbage


Idle length 5



Idle length 20



Idle length 50



Also

- This works for things other than []byte
 - Can be done with arbitrary types
 - Just need some way to reset
- There's a proposal to add something like this to the Go package library
 - sync.Cache/sync.Pool
 - Follow <https://code.google.com/p/go/issues/detail?id=4720>