

CloudFlare's Lua WAF

John Graham-Cumming
October 2014

Two Things

ngx_lua/OpenResty Rules


Two Things

ngx_lua/OpenResty Rules

LuaJIT Rules

Sorry, you have been blocked

You are unable to access jgc.org





OWASP

Open Web Application
Security Project

OWASP ModSecurity Core Rule Set (CRS)

The OWASP ModSecurity CRS Project's goal is to provide an easily "pluggable" set of generic attack detection rules that provide a base level of protection for any web application.

Introduction

The OWASP ModSecurity CRS is a set of web application defense rules for the open source, cross-platform [ModSecurity](#) Web Application Firewall (WAF).

Description

The OWASP ModSecurity CRS provides protections if the following attack/threat categories:

- HTTP Protection - detecting violations of the HTTP protocol and a locally defined usage policy.
- Real-time Blacklist Lookups - utilizes 3rd Party IP Reputation
- HTTP Denial of Service Protections - defense against HTTP Flooding and Slow HTTP DoS Attacks.
- Common Web Attacks Protection - detecting common web application security attack.
- Automation Detection - Detecting bots, crawlers, scanners and other surface malicious activity.
- Integration with AV Scanning for File Uploads - detects malicious files uploaded through the web application.
- Tracking Sensitive Data - Tracks Credit Card usage and blocks leakages.
- Trojan Protection - Detecting access to Trojans horses.
- Identification of Application Defects - alerts on application misconfigurations.
- Error Detection and Hiding - Disguising error messages sent by the server.



CLOUDFLARE



Drupal 7 SA-CORE-2014-005 SQL Injection Protection

16 Oct 2014 by [John Graham-Cumming](#).

9 Share 2 Like 55 Tweet 15

Yesterday the Drupal Security Team released a [critical security patch](#) for Drupal 7 that fixes a very serious [SQL injection](#) vulnerability. At the same time we pushed an update to our Drupal WAF rules to mitigate this problem. Any customer using the WAF and with the Drupal ruleset enabled will have received automatic protection.

CloudFlare Drupal

This ruleset should only be enabled if the Drupal CMS is used for this domain. It contains additional rules that complement the technology-specific protections provided by similar rules in the OWASP ruleset.

Triggered
0 times

☒ ON

Rules


ID	Description	Triggered	Status
D0002	Block requests with odd array arguments	0 times	Default
D0000	Block Large Requests to xmlrpc.php for Drupal CMS	0 times	Default (Bl...
D0001	Block Requests to xmlrpc.php for Drupal CMS	0 times	Default (Di...

Rule D0002 provides protection against this vulnerability. If you do not have that ruleset enabled and are using Drupal clicking the ON button next to CloudFlare Drupal in the WAF Settings will enable protection immediately.



Patching a WHMCS zero day on day zero

03 Oct 2013 by [Dane Knecht](#).



A critical zero-day vulnerability was published today affecting any hosting provider using WHMCS. As part of building a safer web, CloudFlare has added a ruleset to our Web Application Firewall (WAF) to block the published attack vector. Hosting partners running their WHMCS behind CloudFlare's WAF can enable the WHMCS Ruleset and implement [best practices](#) to be fully protected from the attack.


Our friends at WHMCS quickly published a patch here: <http://blog.whmcs.com/?t=79427>

CloudFlare recommends applying the patch for your current version of WHMCS or updating WHMCS to version 5.2.8 to close this vulnerability.

Rule counts

- 5,682 general rules plus 1,937 string matches
- 102 CloudFlare rules
- Customer specific rules
- BAD NEWS: worst case is all the time (run all rules)


<1ms latency



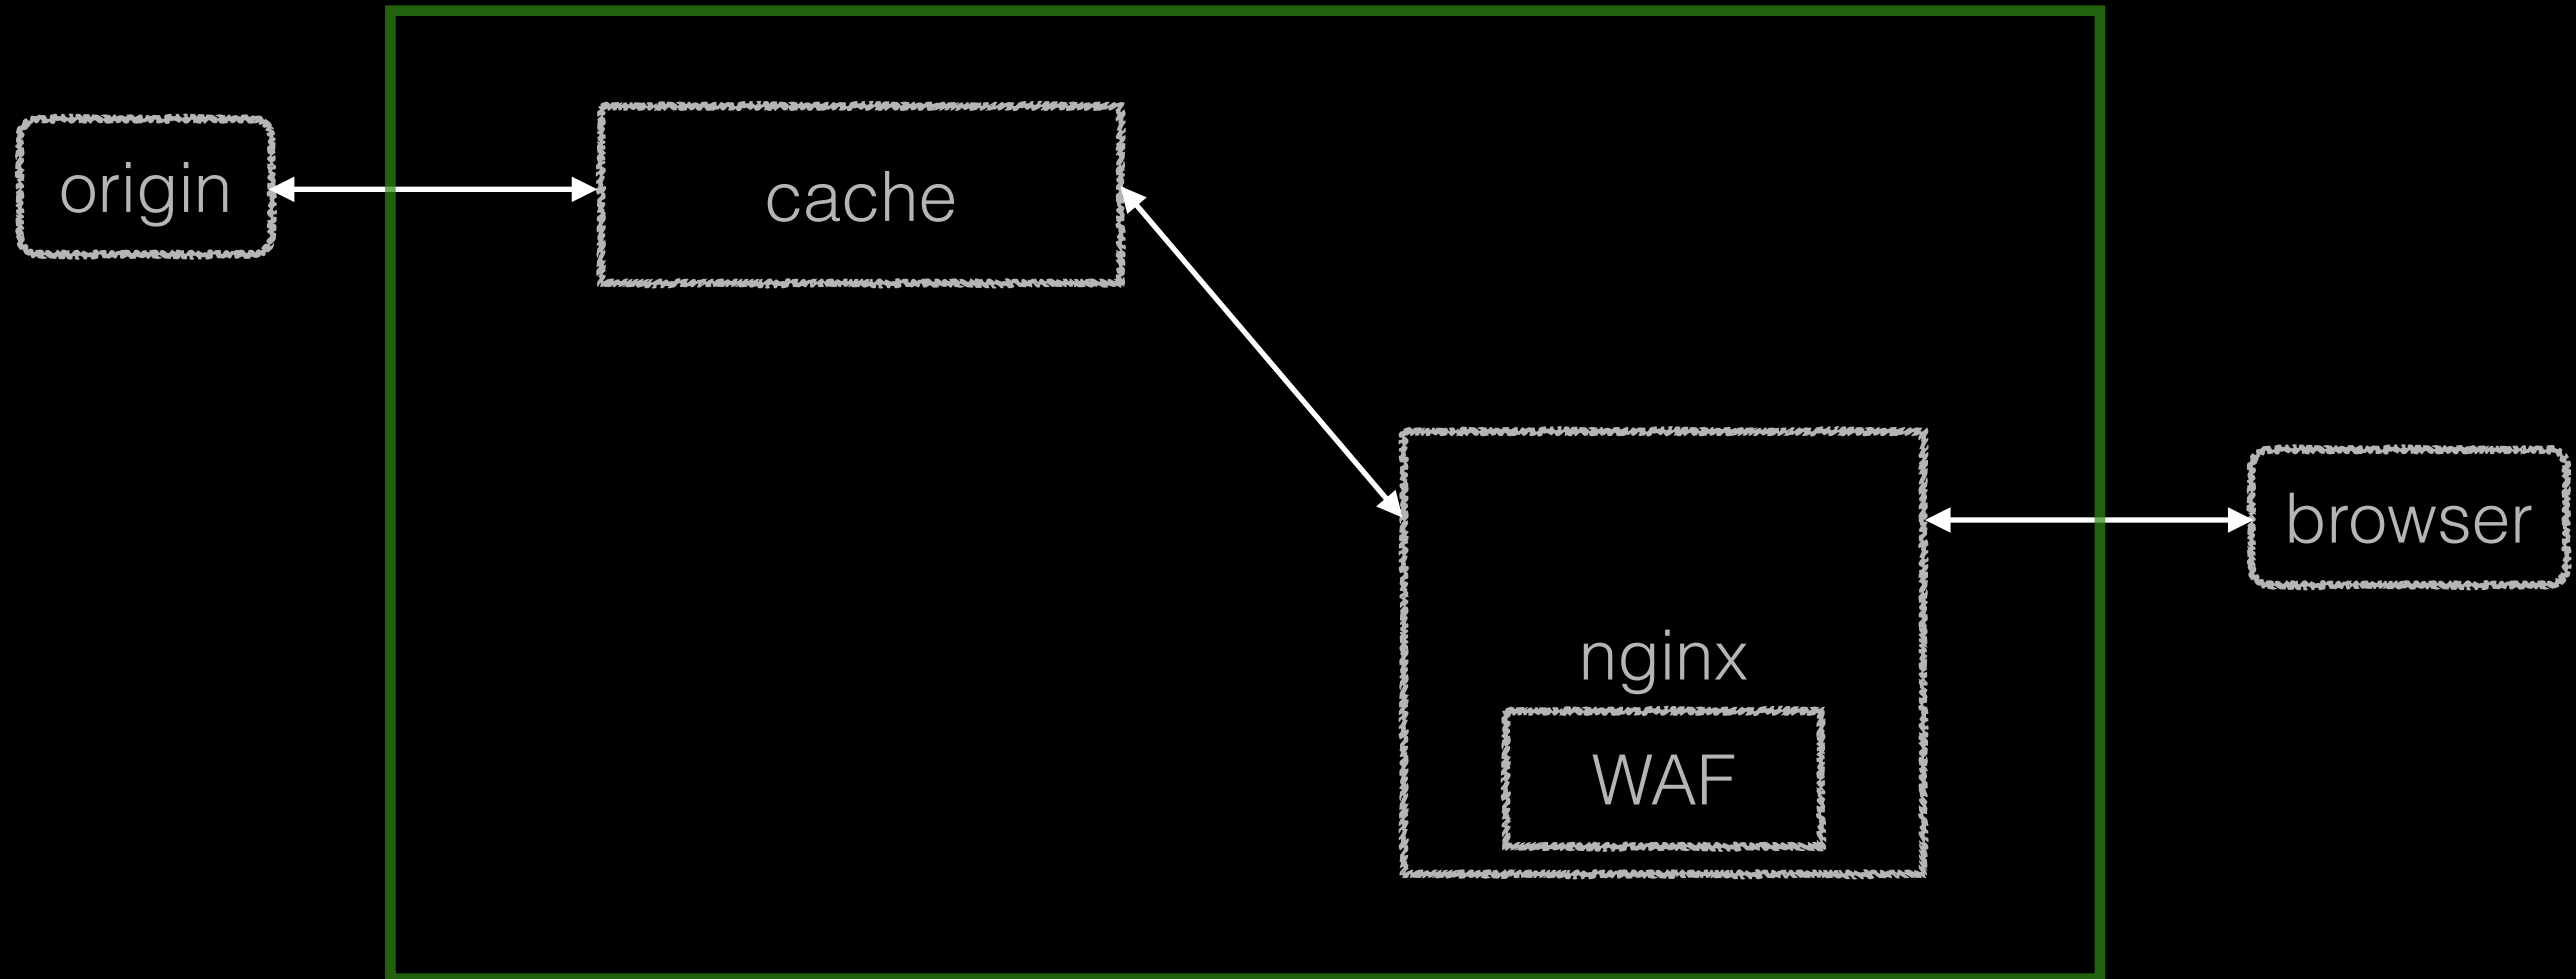
mod_security

- A well established Apache-based project to block bad HTTP requests (and responses)
- Didn't scale for our needs (and we use nginx)
- I essentially rewrote it in Lua and it runs inside nginx
- agentzh did a lot of optimization work (*talk at 1145*)
- It's very, very fast; it runs existing mod_security rules

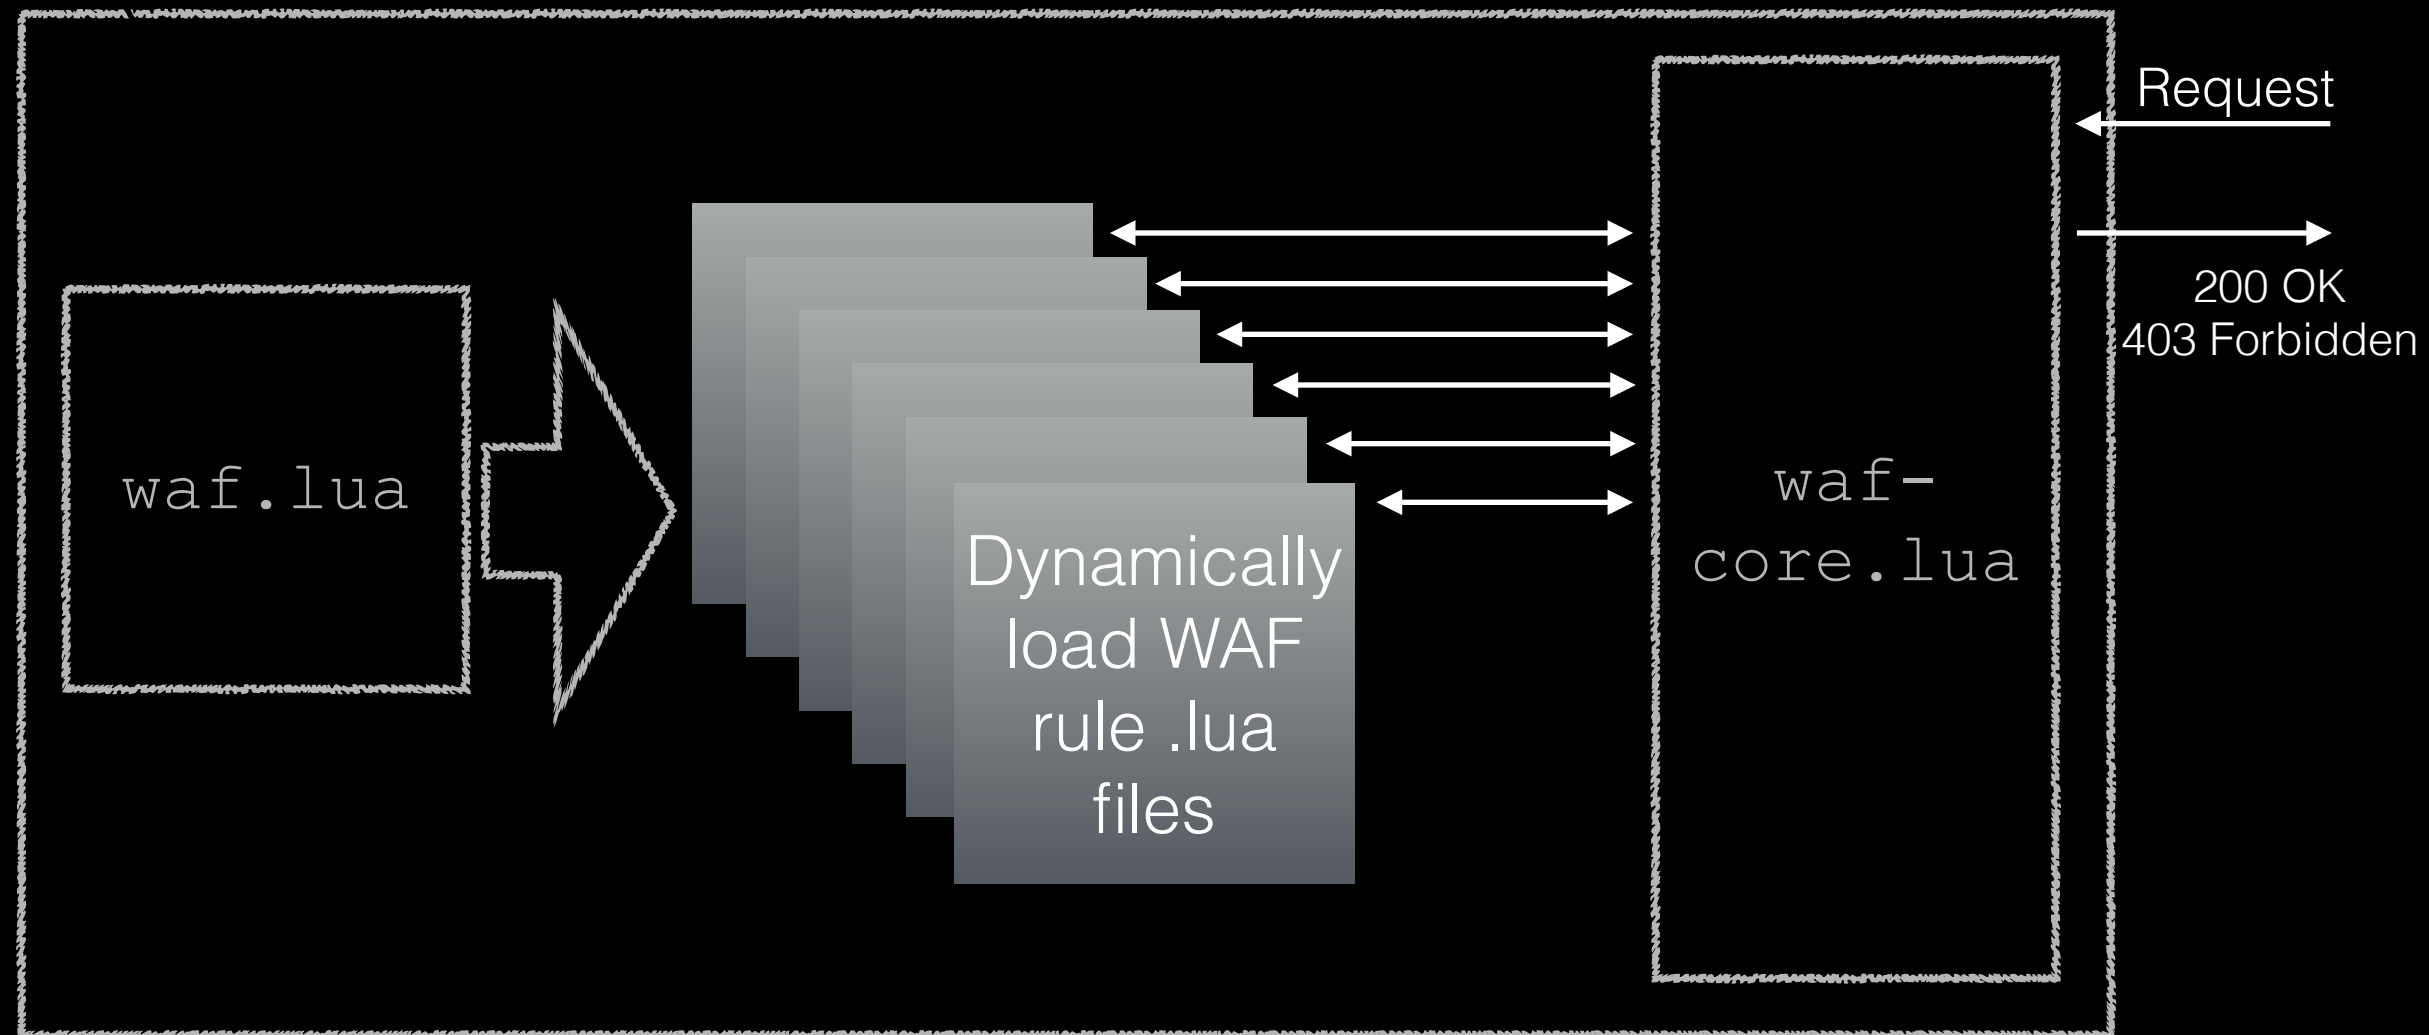
original architecture



new architecture



Inside WAF



nginx.conf

```
location / {  
    set $backend_waf      "WAF_CORE";  
    default_type          'text/plain';  
  
    access_by_lua '        '  
        local waf = require "waf"  
        waf.execute("")  
    ';  
  
    log_by_lua_file "lua/metrics/waf_metrics_main.lua";  
  
    content_by_lua 'ngx.say("")';  
    error_page 500 =200 @error;  
}
```



CLOUDFLARE

Lifecycle of a rule #1

Start with a mod_security rule like this:

```
SecRule ARGS:comment_post_id "@rx ^(\d+)$"  
"drop,id:WP0005,msg:'Exploit DB 28485 Blind  
SQL Injection',phase:0"
```

```
SecRule VARIABLES FUNCTION ACTIONS/  
TRANSFORMATIONS
```



Or our own language

```
SecRule REQUEST_HEADERS:User-Agent "@beginsWith DataStore/"  
    "id:100000,phase:0,t:none,deny,chain,msg:'DataStore Attack'"  
SecRule REQUEST_METHOD "@streq GET" "chain"  
    SecRule REQUEST_URI "\/\?-?\d+=-?\d+" ""
```

```
rule 100000 DataStore Attack  
    REQUEST_HEADERS:User-Agent has-prefix DataStore/ and  
    REQUEST_METHOD is GET and  
    REQUEST_URI matches \/\?-?\d+=-?\d+  
    deny
```



CLOUDFLARE

Lifecycle of a rule #2

It becomes a .json object

```
[{"source": "rulesets/cloudflare/enabled_rules/cloudflare_01_wordpress.conf:60",
  "actions": [
    {"action": "drop"},
    {"parameter": "WP0005", "action": "id"},
    {"parameter": "Exploit DB 28485 Blind SQL Injection", "action": "msg"},
    {"parameter": "0", "action": "phase"}],
  "default": "BLK",
  "variables": ["ARGS:comment_post_id"],
  "description": "Exploit DB 28485 Blind SQL Injection",
  "operator": "^(/d+)$",
  "overrides": ["DEF", "DIS", "SIM", "BLK", "CHL"],
  "directive": "rule",
  "id": "WP0005"}]
```



CLOUDFLARE

Lifecycle of a rule #3

And it turns into Lua code that can run in the WAF

```
if not waf_disabled_ids['WP0005'] and waf_rx(waf, _M.v151_9, '151_9',  
t151_1, '151_1', '^(/d+)$', false, nil, true) then  
    waf_vars['RULE']['ID'] = 'WP0005'  
    waf_activate(waf, _M.rulefile)  
    waf_msg(waf, 'Exploit DB 28485 Blind SQL Injection')  
    waf_drop(waf, _M.rulefile, false)  
end
```



Real example

```
local waf_vars = waf.vars
local waf_streq = waf.streq
local waf_setvar = waf.setvar
local waf_msg = waf.msg
local waf_drop = waf.drop
local waf_disabled_ids = waf.disabled_ids
local waf_deny = waf.deny
local waf_activate = waf.activate
local t1_1 = {}
if not waf_disabled_ids['00001'] and waf_streq(waf, v2_5, '2_5', t1_1, '1_1', 'b783efc191a7c066c1d87068f63a84a39f9830bb', false) then
  waf_vars['RULE']['ID'] = '00001'
  waf_activate(waf, rulefile)
  waf_msg(waf, 'CloudFlare Test Rule (drop) activated')
  waf_setvar(waf, {'TX:ANOMALY_SCORE', '+100'}, {'TX:%{RULE:ID}', 'CloudFlare unique hash test rule (drop)'})
  waf_drop(waf, rulefile)
end
if not waf_disabled_ids['00002'] and waf_streq(waf, v2_5, '2_5', t1_1, '1_1', '4709edce126971876b547523778fa7b942ec14b5', false) then
  waf_vars['RULE']['ID'] = '00002'
  waf_activate(waf, rulefile)
  waf_msg(waf, 'CloudFlare Test Rule (deny) activated')
  waf_setvar(waf, {'TX:ANOMALY_SCORE', '+100'}, {'TX:%{RULE:ID}', 'CloudFlare unique hash test rule (deny)'})
  waf_deny(waf, rulefile)
end
```



Available variables

- Entire mod_security language works with our WAF
- ARGV, ARGV_GET, ARGV_POST, ARGV_NAME, ARGV_COMBINED_SIZE
AUTH_TYPE, QUERY_STRING, REMOTE_USER
REQUEST_COOKIES, REQUEST_COOKIES_NAMES
REQUEST_LINE, REQUEST_METHOD,
REQUEST_PROTOCOL, REQUEST_URI
REQUEST_HEADERS
- REQUEST_HEADERS:User-Agent, ARGV:user_id



Matching functions

- All the mod_security matching functions
- Strings: `@rx` (PCRE regular expression), `@contains`, `@containsWord`, `@beginsWith`, `@endsWith`, `@pm` (match against multiple patterns), `@streq` (string equal), `@empty`, `@notEmpty`, `@hasNull`, `@within`
- Numbers: `@ge`, `@le`, `@gt`, `@lt`, `@eq`, `@ne`
- Encodings: `@validateUrlEncoding`, `@validateUtf8Encoding`, `@validateByteRange`
- `@verifyCC` (check credit card number)

Transformations

- lowercase, trim, trimLeft, trimRight, replaceComments, removeNulls, replaceNulls, removeWhitespace, compressWhitespace, hexEncode, hexDecode, base64Encode, base64Decode,
- urlDecode, urlEncode, cssDecode, jsDecode, escapeSeqDecode,
- md5, sha1, length
- normalisePath, normalisePathWindows,

Web Application Firewall

cloudflare.com

CloudFlare's Web Application Firewall stops real-time attacks like SQL injection, cross-site scripting (XSS), comment spam and other abuse at the network edge. Default settings include coverage for the OWASP core vulnerabilities. You may enable or disable individual rules below.

WAF Rules

WAF Events

Rule Packages

Package Name	Description	Threshold	Action
<input type="checkbox"/> Default Package	Built with OWASP Rules.	Low	Challenge
<input type="checkbox"/> <i>Bad Robots</i>	Detection of bad web robots that are not from search engines but perform malicious searching and spidering of web sites.	Triggered 10 times	<input checked="" type="checkbox"/> ON
<input type="checkbox"/> <i>Generic Attacks</i>	Detection of generic attacks against web-based applications without specific knowledge of the application. Detects things like attempting to access an LDAP directory, inject shell commands, and attacks against PHP.	Triggered 8 times	<input checked="" type="checkbox"/> ON
Rules			
ID	Description	Triggered	Status
950000	Session Fixation	2 times	<input checked="" type="checkbox"/> ON
950002	System Command Access	0 times	<input checked="" type="checkbox"/> ON

Rule compiler optimizations

- Clause reordering so that rules can be quickly skipped if a sub-clause doesn't match
- Regular expression optimization and simplification
- Operator replacement so that fast operators (such as simple string matches) are used where possible
- Providing hints to the WAF runtime about whether macro expansion is needed.
- Global optimizations such as recognizing repeated use of the same strings or variables and ensuring that they are computed only once

Rule compiler optimizations

Swapping `rx(\bactivexobject\b)` for `containsWord(activexobject)`

Swapping `rx(,)` for `contains(,)`

Swapping `rx(^$)` for `empty`

Swapping `rx(^(.*)$)` for `unconditionalMatch`

Swapping `rx(^GET /$)` for `streq(GET /)`

Swapping `rx(Mozilla\5\.0 \ (compatible; MSIE 9\.0; Windows NT 6\.1; WOW64; Trident\5\.0; SLCC2; Media Center PC 6\.0; InfoPath\.3; MS-RTC LM 8; Zune 4\.7\))` for `contains(Mozilla/5.0 (compatible; MSIE 9.0; Windows NT 6.1; WOW64; Trident/5.0; SLCC2; Media Center PC 6.0; InfoPath.3; MS-RTC LM 8; Zune 4.7))`



CLOUDFLARE

Other Optimization

- Basic Lua optimizations
- Deep dive into regular expressions
- Special array handling
- Make sure code is JITable
- Caching
- Memoization

Basic Optimization

Wait til you've finished; Measure; Fix the slow things

Locals way faster than globals

```
local rand = math.random
```

```
local len = #t  
for i=1,len do  
    ...  
end
```

. syntax faster than :

```
local slen = string.len  
s:len() vs. slen(s)
```

Minimize closures




Basic Optimization

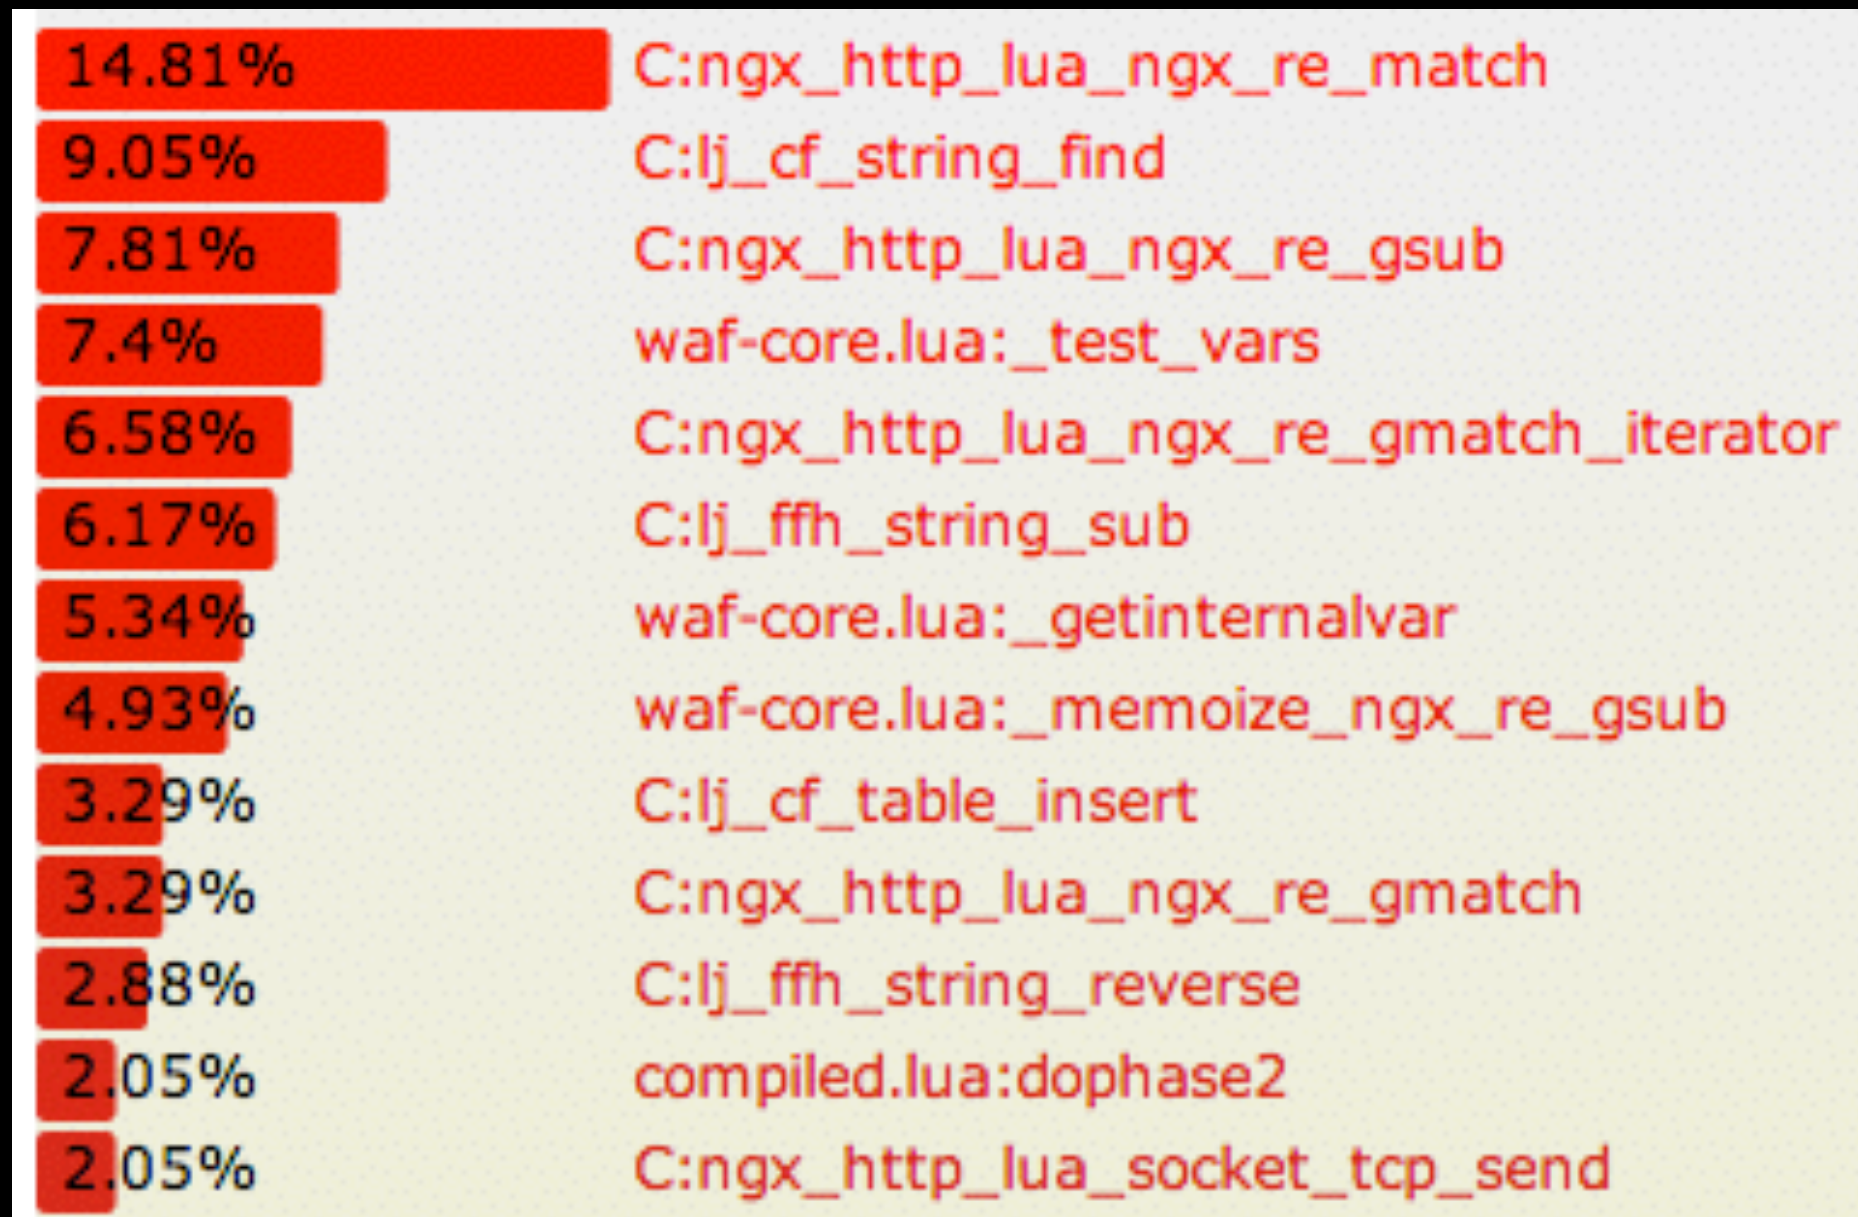
Use PCRE JIT and caching options

Use `lua_shared_dict` to cache frequently accessed items

Lua Flamegraph



Top functions



C Flamegraph



PCRE monitoring

```
$ ./ngx-pcre-stats -p 24528 --exec-time-dist
```

```
Tracing 24528 (/path/to/nginx/sbin/nginx)...
```

```
Hit Ctrl-C to end.
```

```
^C
```

```
Logarithmic histogram for pcre_exec running time distribution  
(us):
```

value	-----	count
0		0
1		0
2	@@@	981
4	@@@	1479
8		16
16		18
32		1

Regex dive

```
$ ./ngx-pcre-stats -p 24528 --total-time-top --luajit20
Tracing 24528 (/path/to/nginx/sbin/nginx)...
Hit Ctrl-C to end.
^C
Top N regexes with longest total running time:
1. pattern /WEB_ATTACK/: 15103us (total data size: 82184)
2. pattern /__cf__\d+/: 11143us (total data size: 25916)
3. pattern /[^\x01-\xff]/: 10233us (total data size:
102825)
4. pattern /\b(?:coalesce\b|root\@)/: 7017us (total data
size: 78230)
5. pattern /(Content-Length|Transfer-Encoding)/: 6766us
(total data size: 17871)
...
```



Slow regexps

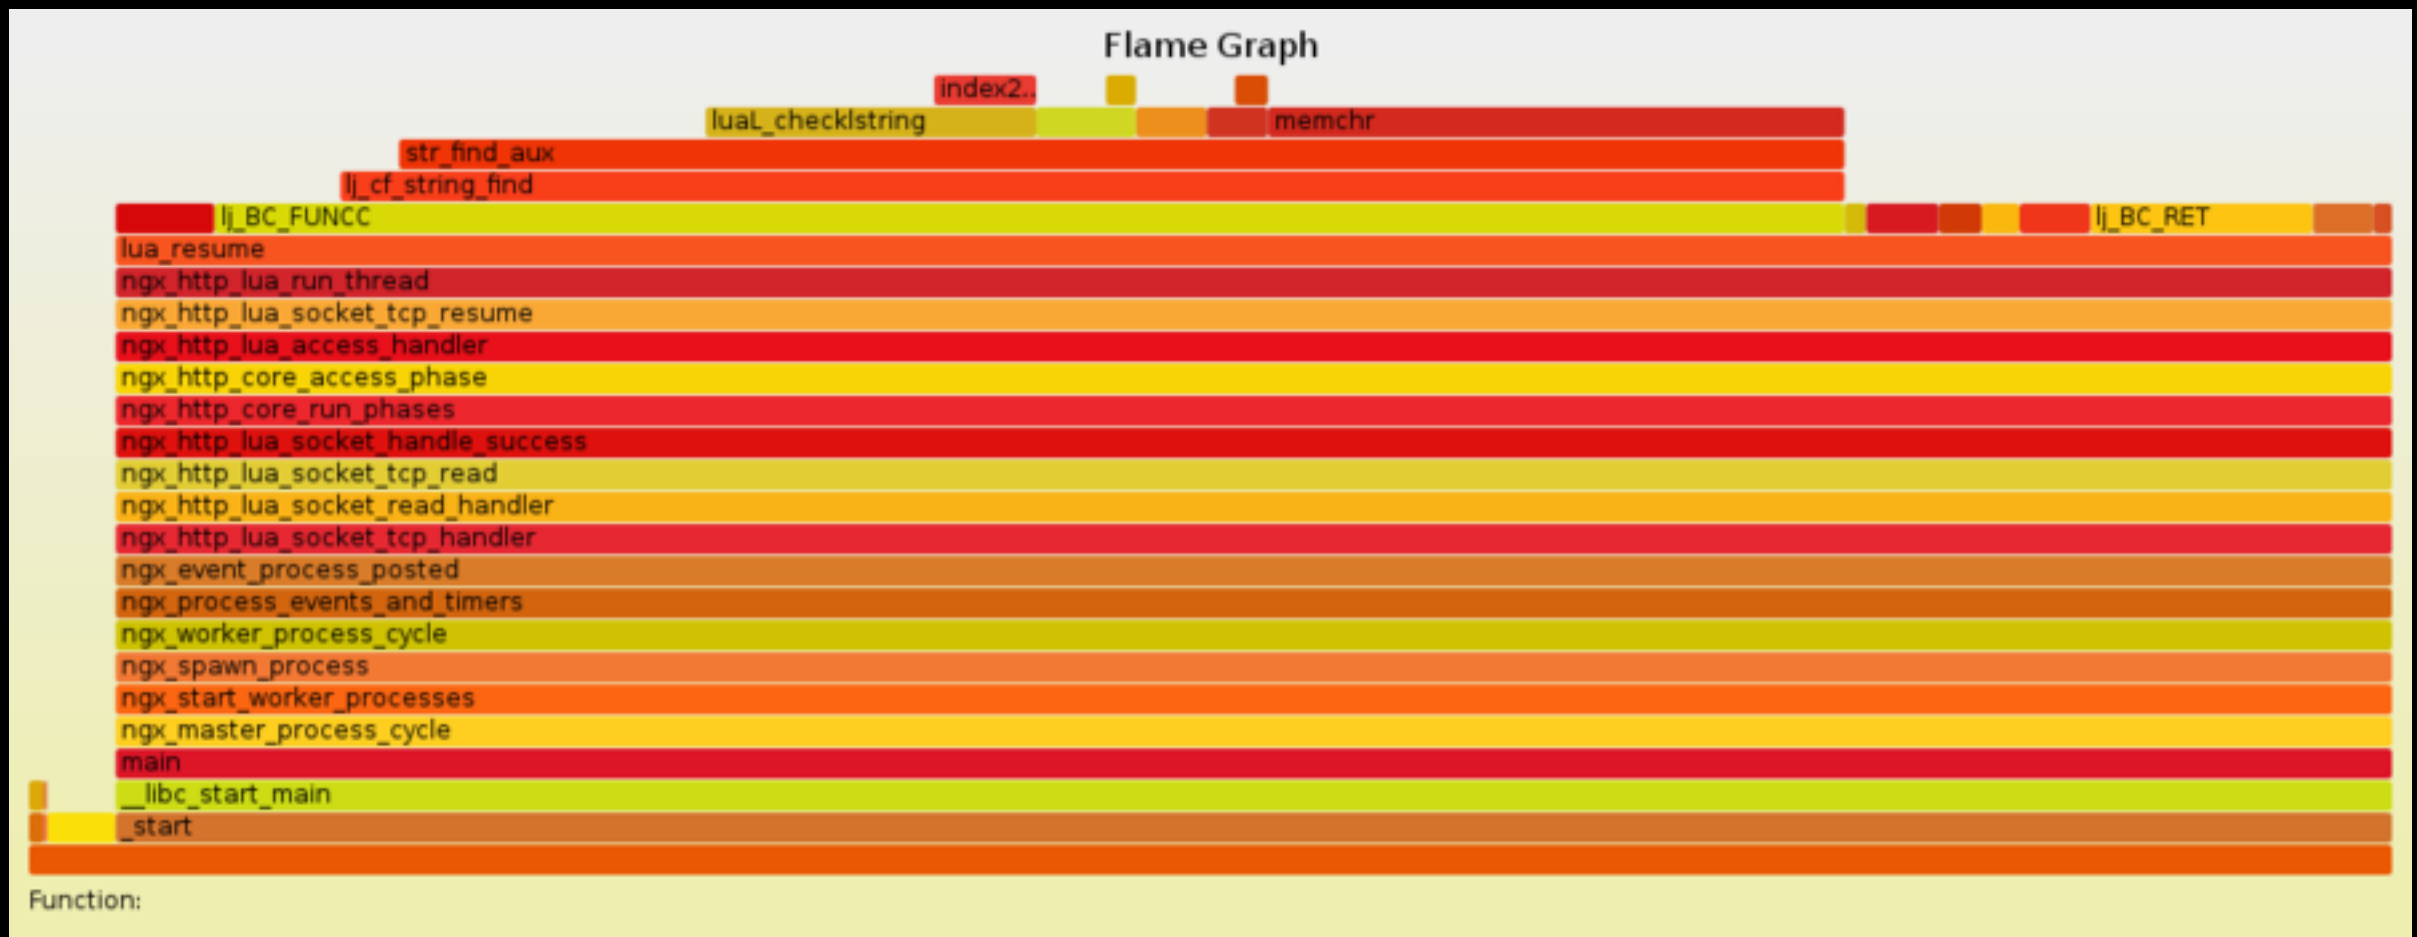
```
$ ./ngx-pcre-stats -p 24528 --worst-time-top --luajit20
Tracing 24528 (/path/to/nginx/sbin/nginx)...
Hit Ctrl-C to end.
^C
```

Top N regexes with worst running time:

1. pattern /\.cookie\b.*?\;\W*?domain\W*?\=/: 98us (data size: 36)
2. pattern /(Content-Length|Transfer-Encoding)/: 89us (data size: 14)
3. pattern /__cf__\d+/: 63us (data size: 8)
4. pattern /[^\x01-\xff]/: 53us (data size: 13)
5. pattern /\b(background|dynsrc|href|lowsrc|src)\b\W*?=/: 53us (data size: 5147)
6. pattern /(?i:<embed[/+\t].*?SRC.*?=/: 47us (data size: 304)
7. pattern /(fromcharcode|alert|eval)\s*\(/: 45us (data size: 24)
8. pattern /\bselect\b.*?\bto_number\b/: 40us (data size: 5147)



Something's spinning



Peek inside backtrace

```
$ ./ngx-lua-bt -p 7599 --luajit20
WARNING: Tracing 7599 (/path/to/nginx/sbin/nginx) for
LuaJIT 2.0...
C:lj_cf_string_find
@/waf/lua/waf-core.lua:201
@/waf/lua/waf-core.lua:676
@/waf/lua/waf-core.lua:1467
@/waf/lua/waf-core.lua:1074
@/waf/lua/rules/oldwaf/compiled.lua:371
@/waf/lua/waf.lua:57
C:lj_ffh_pcall
@/waf/lua/waf.lua:50
access_by_lua:1
```



CLOUDFLARE

Tools

<https://github.com/openresty/nginx-systemtap-toolkit/>

<https://github.com/brendangregg/FlameGraph>

Two step caching

- Compiled Lua stored in memcached
- Once loaded stored in `lua_shared_dict`

Cached require

```
if not package.loaded[r] then
    local src, flags, err = memc:get(filekey)
    local m, err = loadstring(src, filekey)

    if m then
        package.loaded[r] = m()
    end
end

local loaded = require(r)
if not loaded then
    package.loaded[r] = nil
    return nil, "require(" .. r .. ") returned nil"
else
    return loaded, nil
end
```



LuaJIT

- Is a *tracing* compiler for Lua
- Fast!
- Key is to make sure it can JIT
- Has special APIs
- Sponsor Mike Pall!

LuaJIT table API

- `tn = require("table.new")`
`t = tn(narr, nrec)`

Dual use arrays

- Store both array style and hash style in same array
- Get fast iteration of keys or values

```
days[0] = 0
days[1] = "January"
days[2] = 31
days[days[1]] = days[2]
days[0] = days[0] + 2
...
```

```
for i = 1, days[0], 2 do
    k = days[i]
    v = days[i+1]
end
```

```
function add(t, m, d)
    t[m] = d
    t[t[0]+1] = m
    t[t[0]+2] = d
    t[0] = t[0] + 2
end

local tn = require("table.new")
local t = tn(10, 10)
t[0] = 0

add(t, "January", 31)
add(t, "February", 28)
add(t, "March", 31)
add(t, "April", 30)
add(t, "May", 31)
add(t, "June", 30)
add(t, "July", 31)
add(t, "August", 31)
add(t, "September", 30)
add(t, "October", 31)
add(t, "November", 31)
add(t, "December", 31)

for i = 1, t[0], 2 do
    print(t[i], " has ", t[i+1], " days")
end

print("October has ", t["October"], " days")
```



```

% luajit -bl test.lua
-- BYTECODE -- test.lua:1-6
0001      TSETV      2      0      1
0002      TGETB      3      0      0
0003      ADDVN      3      3      0      ; 1
0004      TSETV      1      0      3
0005      TGETB      3      0      0
0006      ADDVN      3      3      1      ; 2
0007      TSETV      2      0      3
0008      TGETB      3      0      0
0009      ADDVN      3      3      1      ; 2
0010      TSETB      3      0      0
0011      RET0       0      1

-- BYTECODE -- test.lua:0-31
0001      FNEW       0      0      ; test.lua:1
0002      GSET       0      1      ; "add"
0003      GGET       0      2      ; "require"
0004      KSTR       1      3      ; "table.new"
0005      CALL       0      2      2
0006      MOV        1      0
0007      KSHORT     2      10
0008      KSHORT     3      10
0009      CALL       1      2      3
0010      KSHORT     2      0
0011      TSETB      2      1      0
0012      GGET       2      1      ; "add"
0013      MOV        3      1
0014      KSTR       4      4      ; "January"
0015      KSHORT     5      31
0016      CALL       2      1      4
0017      GGET       2      1      ; "add"
0018      MOV        3      1
0019      KSTR       4      5      ; "February"
0020      KSHORT     5      28
0021      CALL       2      1      4
0022      GGET       2      1      ; "add"
0023      MOV        3      1
0024      KSTR       4      6      ; "March"
0025      KSHORT     5      31
0026      CALL       2      1      4
0027      GGET       2      1      ; "add"

```


LuaJIT 2.1 now JITable

- CAT (concatenate operator)
- FUNC, FUNCC (C function calls)
- TSETM (return table with multiple values)
- string.char, string.find, string.format, string.lower, string.rep, string.reverse, string.upper
- table.concat, table.foreachi, table.remove

LuaJIT FFI

```
local ffi = require 'ffi'
local C = ffi.C

ffi.cdef[[
    typedef long time_t;

    typedef struct timeval {
        time_t tv_sec;
        time_t tv_usec;
    } timeval;

    int gettimeofday(struct timeval* t, void* tzp);
]]

local gettimeofday_struct = ffi.new("timeval")

-- _gettimeofday_: wrapper function that calls the C gettimeofday
-- function via FFI and returns a value in microseconds
local function gettimeofday()
    C.gettimeofday(gettimeofday_struct, nil)
    return tonumber(gettimeofday_struct.tv_sec) * 1000000 +
        tonumber(gettimeofday_struct.tv_usec)
end
```



CLOUDFLARE

Closing Thoughts

ngx_lua/OpenResty Rules

LuaJIT Rules

PCRE Rules