



Lua: the world's most infuriating language

October 17, 2013

The simplest way to a safer, faster and smarter website

My background

- Lots of...
 - Assembly
 - C, C++ and Go
 - Perl, Tcl
 - LISP (and relations)
- And lately...
 - Lua







CHDK on Canon A560



CHDK Lua Interface

- Lua's extensibility means it can control the camera
- <http://chdk.wikia.com/wiki/Lua>
- Functions added to Lua like
 - `shoot()` – takes a picture
 - `press()`, `release()` – press the shutter button (allows a partial press)
 - `get_orientation_sensor()` – figure out camera orientation
 - `get_temperature()` – get camera internal temperatures

```
-- Now enter a self-check of the manual mode settings

log( "Self-check started" )

assert_prop( 49, -32764, "Not in manual mode" )
assert_prop( 5, 0, "AF Assist Beam should be Off" )
assert_prop( 6, 0, "Focus Mode should be Normal" )
assert_prop( 8, 0, "AiAF Mode should be On" )
assert_prop( 21, 0, "Auto Rotate should be Off" )
assert_prop( 29, 0, "Bracket Mode should be None" )
assert_prop( 57, 0, "Picture Mode should be Superfine" )
assert_prop( 66, 0, "Date Stamp should be Off" )
assert_prop( 95, 0, "Digital Zoom should be None" )
assert_prop( 102, 0, "Drive Mode should be Single" )
assert_prop( 133, 0, "Manual Focus Mode should be Off" )
assert_prop( 143, 2, "Flash Mode should be Off" )
assert_prop( 149, 100, "ISO Mode should be 100" )
assert_prop( 218, 0, "Picture Size should be L" )
assert_prop( 268, 0, "White Balance Mode should be Auto" )
assert_gt( get_time("Y"), 2009, "Unexpected year" )
assert_gt( get_time("h"), 6, "Hour appears too early" )
assert_lt( get_time("h"), 20, "Hour appears too late" )
assert_gt( get_vbatt(), 3000, "Batteries seem low" )
assert_gt( get_jpg_count(), ns, "Insufficient card space" )

log( "Self-check complete" )
```


<https://github.com/jgrahamc/gaga/tree/master/gaga-1/camera>

```
if ( ok == 1 ) then
  sleep(s)
  log( "Starting picture capture" )

  n = 0

  while ( 1 ) do
    tc = c
    while ( tc > 0 ) do
      shoot()
      n = n + 1
      log( string.format("Picture %i taken", n ))
      tc = tc - 1
    end
    log( string.format("Temperatures: %i, %i, %i",
      get_temperature(0), get_temperature(1), get_temperature(2) ))
    log( string.format("Battery level %i", get_vbatt()))
    sleep(i)
  end
end

log( "Done" )
```



Initial Irritants

- Seemed awfully verbose (bit like BASIC!)
 - `if... then... else... end`
- Little shortcuts were missing
 - `x += 1`
 - `p?a:b`
- Made it feel like a toy language

Irritant: +=

- Happy not to have $x++$
- But why no +=?

Irritant: ternary operator

- Super useful to be able to do

```
local bird = duck?"It's a duck":"Not a duck"
```

- Can't.
- Solution is the ugly

```
local bird = duck and "It's a duck" or "Not a duck"
```

- And it's a bad solution

```
local thing = question and returnsfalse() or crash()
```

Irritant: `not` and `~`

- Not equals is `~`
 - Which Perl programmers will confuse with `=~`

```
if x ~= 5 then print "Not five" end
```

- But `not` is not `~` it's `not`!

✖ `if not x == 5 then print "Not five" end`

✖ `if ~(x == 5) then print "Not five" end`


```
if not (x == 5) then print "Not five" end
```

Next Lua Program


- Web Application Firewall
- About 2,000 lines of Lua
 - Replaced 37,000 line C program!
 - Plus 3,000 of automatically generated Lua
- Used by CloudFlare to process HTTP requests
 - Checks for XSS, CSRF, SQL injection
 - Bad browsers
 - Custom rules
- Turns out Lua was great for this!



Really, really fast: 1 to 2ms per request



LuaJIT



LuaJIT FFI

```
local ffi = require 'ffi'
local C = ffi.C

ffi.cdef[[
    typedef long time_t;

    typedef struct timeval {
        time_t tv_sec;
        time_t tv_usec;
    } timeval;

    int gettimeofday(struct timeval* t, void* tzp);
]]

local gettimeofday_struct = ffi.new("timeval")

-- _gettimeofday_: wrapper function that calls the C gettimeofday
-- function via FFI and returns a value in microseconds
local function gettimeofday()
    C.gettimeofday(gettimeofday_struct, nil)
    return tonumber(gettimeofday_struct.tv_sec) * 1000000 + tonumber(gettimeofday_struct.tv_usec)
end
```


Lulip

- Line level profiler for Lua in Lua: <https://github.com/jgrahamc/lulip/>

```
local profiler = require 'lulip'  
local p = profiler:new()  
p:dont('some-module')  
p:maxrows(25)  
p:start()
```

```
-- execute code here
```

```
p:stop()  
p:dump(output_file_name)
```

Lulip Output

file:line	count	elapsed (ms)	line
wr.lua:1129	2	822.455	hash = ngx_shal_bin(value)
wr.lua:1172	428	470.849	captures, err = ngx_re_match(v, p)
wr.lua:1197	3762	207.487	x = string_find(v, f)
wr.lua:212	157	154.386	string_gsub(v, "//([^/]+)/", "%1")
wr.lua:1196	3788	87.475	for i=1,g() do
wr.lua:1158	1563	52.906	if not f() then

Lulip Core

- Nice things
 - Lua's debug library
 - LuaJIT's FFI for `gettimeofday()`
 - Closures

```
-- start: begin profiling
function start(self)
    self:dont('lulip.lua')
    self.start_time = gettimeofday()
    self.current_line = nil
    self.current_start = 0

    debug.sethook(function(e,l) self:event(e, l) end, "l")
end
```



Performance Tricks

- Wait til you've finished; Measure; Fix the slow things

- Locals way faster than globals

```
local rand = math.random
```

```
local len = #t
for i=1,len do
    ...
end
```

- . syntax faster than :

```
local slen = string.len
s:len() vs. slen(s)
```

- Minimize closures



Autogenerated Code

```
local waf_vars = waf.vars
local waf_streq = waf.streq
local waf_setvar = waf.setvar
local waf_msg = waf.msg
local waf_drop = waf.drop
local waf_disabled_ids = waf.disabled_ids
local waf_deny = waf.deny
local waf_activate = waf.activate
local t1_1 = {}
if not waf_disabled_ids['00001'] and waf_streq(waf, v2_5, '2_5', t1_1, '1_1', 'b783efc191a7c066c1d87068f63a84a39f9830bb', false) then
  waf_vars['RULE']['ID'] = '00001'
  waf_activate(waf, rulefile)
  waf_msg(waf, 'CloudFlare Test Rule (drop) activated')
  waf_setvar(waf, {{'TX:ANOMALY_SCORE', '+100'},{'TX:%{RULE:ID}', 'CloudFlare unique hash test rule (drop)'}})
  waf_drop(waf, rulefile)
end
if not waf_disabled_ids['00002'] and waf_streq(waf, v2_5, '2_5', t1_1, '1_1', '4709edce126971876b547523778fa7b942ec14b5', false) then
  waf_vars['RULE']['ID'] = '00002'
  waf_activate(waf, rulefile)
  waf_msg(waf, 'CloudFlare Test Rule (deny) activated')
  waf_setvar(waf, {{'TX:ANOMALY_SCORE', '+100'},{'TX:%{RULE:ID}', 'CloudFlare unique hash test rule (deny)'}})
  waf_deny(waf, rulefile)
end
```


nginx + Lua and OpenResty

- <http://wiki.nginx.org/HttpLuaModule>
- Control over all phases of nginx inside Lua
- Very fast, very flexible, non-blocking (without callbacks!)
- Sockets, coroutines, subrequests, shared in-memory caching
- <http://openresty.org/>
- Complete package of nginx, Lua and a set of libraries
- Access to Redis, memcached, MySQL, Postgresql
- JSON parsing, DNS, locks, ...
- CloudFlare HTTP almost entirely nginx + Lua / OpenResty

CloudFlare WAF

```
location / {  
    set $backend_waf      "WAF_CORE";  
    default_type          'text/plain';  
  
    access_by_lua '        local waf = require "waf"  
                           waf.execute("")'  
    ;  
  
    log_by_lua_file "lua/metrics/waf_metrics_main.lua";  
  
    content_by_lua 'ngx.say("")';  
    error_page 500 =200 @error;  
}
```

Irritant: Arrays

- Index starts at 1
- # doesn't do what you think

```
> x = {"a", "b", "c"}
```

```
> =#x
```

```
3
```

```
> x = {"a", "b", nil, "c"}
```

```
> =#x
```

```
4
```

```
> x = {"a", "b", nil, "c", nil}
```

```
> =#x
```

```
2
```

Irritant: escape characters

- `string.find`, `string.match`
- Lots of special characters just like regular expressions
 - Character classes: `[a-z]`, `[^a-z]`, `.`
 - Repetition: `*`, `+`, `?`
 - Anchors: `^` and `$`
 - Groups and alternation: `(foo|bar)`
- And then... `%`
 - `%d`, `%u`, `%D`, `%s`
- Although `%b` and `%f` are cool

Awesome: tables

- Anything can be a key; anything can be a value
 - Tables as values for nesting
 - Functions as values

```
x = {"a", "b", "c"}
```

```
y = {subtable=x, double=function(v) return v*2 end}
```

- Tables are references

```
function f(t, v) t.subtable = v end
```

```
f(y, {"1", "2", "3"})
```


Tables aren't cool... metatables are cool



metatables to make a table read-only

```
local t = {a = "A"}  
local _t = t  
  
t = {}  
local ro = {  
    __index = _t,  
    __newindex = function() error("R/O") end  
}  
  
setmetatable(t, ro)
```

```
print(t.a)
```

✖ t.a = "B"

metatables for lazy loading

```
local t = {}  
local loader = {  
    __index=function(_t, _v)  
        _t[_v] = docostlyload(_v)  
        return _t[_v]  
    end}  
setmetatable(t, loader)  
  
print(t.expensive)
```

metatables to make objects

```
local C = {}  
C.__index = C  
  
function C.new(d)  
    local newObject = {data=d}  
    return setmetatable(newObject, C)  
end  
  
function C.get_data(self)  
    return self.data  
end  
  
local o = C.new("hello")  
print(o:get_data())
```

metatables to sandbox code #1

```
local env = { print = print }  
local envmeta = { __index={}, __newindex=function() end }  
setmetatable(env, envmeta)
```

```
function run(code)  
    local f = loadstring(code)  
    setfenv(f, env)  
    pcall(f)  
end
```

```
run([[  
    local x = "Hello, World!"  
    print(x)  
    local y = string.len(x)  
]])
```



metatables to sandbox code #2

```
local env = { print = print, string = { len = string.len } }  
local envmeta = { __index={}, __newindex=function() end }  
setmetatable(env, envmeta)
```

```
function run(code)  
    local f = loadstring(code)  
    setfenv(f, env)  
    pcall(f)  
end
```

```
run([[  
    local x = "Hello, World!"  
    print(x)  
    local y = string.len(x)  
]])
```


metatables for undefined variable detection

```
function doubler(x) return X*2 end
print(doubler(2))
t.lua:1: attempt to perform arithmetic on global 'X' (a nil value)
```

```
function doubler(x) return X*2 end
catcher={__index=function(t,v) error(v .. " undefined") end,
          __newindex=function(t,k,v) error(v .. " undefined") end}
setmetatable(_G, catcher)
print(doubler(2))
lua: t.lua:2: Tried to read undefined X
```

Conclusion

- Get past initial irritation!
- Lua is a GREAT language for embedding in large systems
 - Fast
 - Lots of functionality
 - Good standard library
 - Small
 - Extensible both from C and to C